

Data Acquisition with APIs in R

You can download this .qmd file from [here](#). Just hit the Download Raw File button.

Credit to Brianna Heggeseth and Leslie Myint from Macalester College for a few of these descriptions and examples.

Getting data from websites

Option 1: APIs

When we interact with sites like The New York Times, Zillow, and Google, we are accessing their data via a graphical layout (e.g., images, colors, columns) that is easy for humans to read but hard for computers.

An **API** stands for **Application Programming Interface**, and this term describes a general class of tool that allows computers, rather than humans, to interact with an organization's data. How does this work?

- When we use web browsers to navigate the web, our browsers communicate with web servers using a technology called [HTTP](#) or Hypertext Transfer Protocol to get information that is formatted into the display of a web page.
- Programming languages such as R can also use HTTP to communicate with web servers. The easiest way to do this is via [Web APIs](#), or Web Application Programming Interfaces, which focus on transmitting raw data, rather than images, colors, or other appearance-related information that humans interact with when viewing a web page.

A large variety of web APIs provide data accessible to programs written in R (and almost any other programming language!). Almost all reasonably large commercial websites offer APIs. Todd Motto has compiled an expansive list of [Public Web APIs](#) on GitHub, although it's about 3 years old now so it's not a perfect or complete list. Feel free to browse this list to see what data sources are available.

For our purposes of obtaining data, APIs exist where website developers make data nicely packaged for consumption. The language HTTP (hypertext transfer protocol) underlies APIs,

and the R package `httr()` (and now the updated `httr2()`) was written to map closely to HTTP with R. Essentially you send a request to the website (server) where you want data from, and they send a response, which should contain the data (plus other stuff).

The case studies in this document provide a really quick introduction to data acquisition, just to get you started and show you what's possible. For more information, these links can be somewhat helpful:

- <https://www.geeksforgeeks.org/functions-with-r-and-rvest/#>
- <https://nceas.github.io/oss-lessons/data-liberation/intro-webscraping.html>

Wrapper packages

In R, it is easiest to use Web APIs through a **wrapper package**, an R package written specifically for a particular Web API.

- The R development community has already contributed wrapper packages for many large Web APIs (e.g. ZillowR, rtweet, genius, Rspotify, tidycensus, etc.)
- To find a wrapper package, search the web for “R package” and the name of the website. For example:
 - Searching for “R Reddit package” returns [RedditExtractor](#)
 - Searching for “R Weather.com package” returns [weatherData](#)
- [rOpenSci](#) also has a good collection of wrapper packages.

In particular, `tidycensus` is a wrapper package that makes it easy to obtain desired census information for mapping and modeling:


Obtaining raw data from the Census Bureau was that easy! Often we will have to obtain and use a secret API key to access the data, but that's not always necessary with `tidycensus`.

Now we can tidy that data and produce plots and analyses. [Here's a decent place to get more information about the variable codes.](#)


```
# Rename cryptic variables from the census form
sample_acs_data <- sample_acs_data |>
  rename(population = B01003_001E,
         population_moe = B01003_001M,
         median_income = B19013_001E,
         median_income_moe = B19013_001M)

# Plot with geom_sf since our data contains 1 row per census tract
#   with its geometry
ggplot(data = sample_acs_data) +
```


```
geom_sf(aes(fill = median_income), colour = "white", linetype = 2) +  
theme_void()
```



```
# The whole state of MN is overwhelming, so focus on a single county  
sample_acs_data |>  
  filter(str_detect(NAME, "Ramsey")) |>  
  ggplot() +  
    geom_sf(aes(fill = median_income), colour = "white", linetype = 2)
```



```
# Look for relationships between variables with 1 row per tract
as_tibble(sample_acs_data) |>
  ggplot(aes(x = population, y = median_income)) +
    geom_point() +
    geom_smooth(method = "lm")
```



Extra resources:

- **tidycensus**: wrapper package that provides an interface to a few census datasets *with map geometry included!*
 - Full documentation is available at <https://walker-data.com/tidycensus/>
- **censusapi**: wrapper package that offers an interface to all census datasets
 - Full documentation is available at <https://www.hrecht.com/censusapi/>

`get_acs()` is one of the functions that is part of **tidycensus**. Let's explore what's going on behind the scenes with `get_acs()`...

Accessing web APIs directly

Getting a Census API key

Many APIs (and their wrapper packages) require users to obtain a **key** to use their services.

- This lets organizations keep track of what data is being used.
- It also **rate limits** their API and ensures programs don't make too many requests per day/minute/hour. Be aware that most APIs do have rate limits — especially for their free tiers.

Navigate to https://api.census.gov/data/key_signup.html to obtain a Census API key:

- Organization: St. Olaf College
- Email: Your St. Olaf email address

You will get the message:

Your request for a new API key has been successfully submitted. Please check your email. In a few minutes you should receive a message with instructions on how to activate your new key.

Check your email. Copy and paste your key into a new text file:

- (In RStudio) File > New File > Text File (towards the bottom of the menu)
- Save as `census_api_key.txt` in the same folder as this `.qmd`.

You could then read in the key with code like this:

```
myapikey <- readLines("C:/Users/charl/Documents/SDS_264/census_api_key")
```

```
Warning in readLines("C:/Users/charl/Documents/SDS_264/census_api_key"):  
incomplete final line found on  
'C:/Users/charl/Documents/SDS_264/census_api_key'
```

Handling API keys

While this works, the problem is once we start backing up our files to GitHub, your API key will also appear on GitHub, and you want to keep your API key secret. Thus, we might use **environment variables** instead:

One way to store a secret across sessions is with environment variables. Environment variables, or envvars for short, are a cross platform way of passing information to processes. For passing envvars to R, you can list name-value pairs in a file called `.Renviron` in your home directory. The easiest way to edit it is to run:

```
file.edit("~/Renviron")  
  
Sys.setenv(PATH = "path", VAR1 = "value1", VAR2 = "value2")
```

The file looks something like

`PATH = "path" VAR1 = "value1" VAR2 = "value2"` And you can access the values in R using `Sys.getenv()`:

```
Sys.getenv("VAR1")  
#> [1] "value1"
```

Note that `.Renviron` is only processed on startup, so you'll need to restart R to see changes.

Another option is to use `Sys.setenv` and `Sys.getenv`:

```
# I used the first line to store my CENSUS API key in .Renviron  
# after uncommenting - should only need to run one time  
#Sys.setenv(CENSUS_API_KEY = "my personal key")  
my_census_api_key <- Sys.getenv("CENSUS_API_KEY")
```

Navigating API documentation

Navigate to the [Census API user guide](#) and click on the “Example API Queries” tab.

Let's look at the Population Estimates Example and the American Community Survey (ACS) Example. These examples walk us through the steps to incrementally build up a URL to obtain desired data. This URL is known as a web API **request**.

https://api.census.gov/data/2019/acs/acs1?get=NAME,B02015_009E,B02015_009M&for=state:*

- <https://api.census.gov>: This is the **base URL**.
 - [http://](#): The **scheme**, which tells your browser or program how to communicate with the web server. This will typically be either **http:** or **https:**.
 - [api.census.gov](#): The **hostname**, which is a name that identifies the web server that will process the request.
- [data/2019/acs/acs1](#): The **path**, which tells the web server how to get to the desired resource.
 - In the case of the Census API, this locates a desired dataset in a particular year.
 - Other APIs allow search functionality. (e.g., News organizations have article searches.) In these cases, the path locates the search function we would like to call.
- [?get=NAME,B02015_009E,B02015_009M&for=state:*](#): The **query parameters**, which provide the parameters for the function you would like to call.
 - We can view this as a string of key-value pairs separated by **&**. That is, the general structure of this part is **key1=value1&key2=value2**.

key	value
key	value
get	NAME,B02015_009E,B02015_009M
for	state:*

Typically, each of these URL components will be specified in the API documentation. Sometimes, the scheme, hostname, and path (<https://api.census.gov/data/2019/acs/acs1>) will be referred to as the **endpoint** for the API call.

We will first use the [httr2 package](#) to build up a full URL from its parts.

- `request()` creates an API request object using the **base URL**
- `req_url_path_append()` builds up the URL by adding path components separated by `/`
- `req_url_query()` adds the `?` separating the endpoint from the query and sets the key-value pairs in the query
 - The `.multi` argument controls how multiple values for a given key are combined.
 - The `I()` function around `"state:*` inhibits parsing of special characters like `:` and `*`. (It's known as the "as-is" function.)
 - The backticks around `for` are needed because `for` is a reserved word in R (for for-loops). You'll need backticks whenever the key name has special characters (like spaces, dashes).
 - We can see from [here](#) that providing an API key is achieved with `key=YOUR_API_KEY`.

```
# Request total number of Hmong residents and margin of error by state
# in 2019, as in the User Guide
CENSUS_API_KEY <- Sys.getenv("CENSUS_API_KEY")
req <- request("https://api.census.gov") |>
  req_url_path_append("data") |>
  req_url_path_append("2019") |>
  req_url_path_append("acs") |>
  req_url_path_append("acs1") |>
  req_url_query(get = c("NAME", "B02015_009E", "B02015_009M"), `for` = I("state:*"), key =
```

Why would we ever use these steps instead of just using the full URL as a string?

- To generalize this code with functions! (This is exactly what wrapper packages do.)
- To handle special characters

- e.g., query parameters might have spaces, which need to be represented in a particular way in a URL (URLs can't contain spaces)

Once we've fully constructed our request, we can use `req_perform()` to send out the API request and get a **response**.

```
resp <- req_perform(req)
resp
```

We see from **Content-Type** that the format of the response is something called JSON. We can navigate to the request URL to see the structure of this output.

- JSON (Javascript Object Notation) is a nested structure of key-value pairs.
- We can use `resp_body_json()` to parse the JSON into a nicer format.
 - Without `simplifyVector = TRUE`, the JSON is read in as a list.

```
resp_json_list <- resp |> resp_body_json()
head(resp_json_list, 2)
```

```
[[1]]
[[1]][[1]]
[1] "NAME"

[[1]][[2]]
[1] "B02015_009E"

[[1]][[3]]
[1] "B02015_009M"

[[1]][[4]]
[1] "state"

[[2]]
[[2]][[1]]
[1] "Mississippi"

[[2]][[2]]
NULL

[[2]][[3]]
```

NULL

```
[[2]][[4]]  
[1] "28"
```

```
resp_json_df <- resp |> resp_body_json(simplifyVector = TRUE)  
head(resp_json_df)
```

	[,1]	[,2]	[,3]	[,4]
[1,]	"NAME"	"B02015_009E"	"B02015_009M"	"state"
[2,]	"Mississippi"	NA	NA	"28"
[3,]	"Missouri"	"953"	"1141"	"29"
[4,]	"Montana"	NA	NA	"30"
[5,]	"Nebraska"	"412"	"477"	"31"
[6,]	"Nevada"	"863"	"745"	"32"

```
resp_json_df <- janitor::row_to_names(resp_json_df, 1)  
head(resp_json_df)
```

	NAME	B02015_009E	B02015_009M	state
[1,]	"Mississippi"	NA	NA	"28"
[2,]	"Missouri"	"953"	"1141"	"29"
[3,]	"Montana"	NA	NA	"30"
[4,]	"Nebraska"	"412"	"477"	"31"
[5,]	"Nevada"	"863"	"745"	"32"
[6,]	"New Hampshire"	NA	NA	"33"

All right, let's try this! First we'll grab total population and median household income for all census tracts in MN using 3 approaches

```
# First using tidycensus  
library(tidycensus)  
sample_acs_data <- tidycensus::get_acs(  
  year = 2021,  
  state = "MN",  
  geography = "tract",  
  variables = c("B01003_001", "B19013_001"),  
  output = "wide",  
  geometry = TRUE,  
  county = "Hennepin", # specify county in call  
  show_call = TRUE     # see resulting query  
)
```

```
# Next using httr2
req <- request("https://api.census.gov") |>
  req_url_path_append("data") |>
  req_url_path_append("2020") |>
  req_url_path_append("acs") |>
  req_url_path_append("acs5") |>
  req_url_query(get = c("NAME", "B01003_001E", "B19013_001E"), `for` = I("tract:*"), `in` =

resp <- req_perform(req)
resp
resp_json_df <- resp |> resp_body_json(simplifyVector = TRUE)
head(resp_json_df)
```

	[,1]		[,2]
[1,]	"NAME"		"B01003_001E"
[2,]	"Census Tract 1.01, Hennepin County, Minnesota"		"3472"
[3,]	"Census Tract 1.02, Hennepin County, Minnesota"		"4992"
[4,]	"Census Tract 3, Hennepin County, Minnesota"		"3404"
[5,]	"Census Tract 6.01, Hennepin County, Minnesota"		"4706"
[6,]	"Census Tract 6.03, Hennepin County, Minnesota"		"3301"

	[,3]	[,4]	[,5]	[,6]
[1,]	"B19013_001E"	"state"	"county"	"tract"
[2,]	"70927"	"27"	"053"	"000101"
[3,]	"46333"	"27"	"053"	"000102"
[4,]	"82098"	"27"	"053"	"000300"
[5,]	"71122"	"27"	"053"	"000601"
[6,]	"96875"	"27"	"053"	"000603"

```
resp_json_df <- janitor::row_to_names(resp_json_df, 1)
head(resp_json_df)
```


	NAME	B01003_001E	B19013_001E
[1,]	"Census Tract 1.01, Hennepin County, Minnesota"	"3472"	"70927"
[2,]	"Census Tract 1.02, Hennepin County, Minnesota"	"4992"	"46333"
[3,]	"Census Tract 3, Hennepin County, Minnesota"	"3404"	"82098"
[4,]	"Census Tract 6.01, Hennepin County, Minnesota"	"4706"	"71122"
[5,]	"Census Tract 6.03, Hennepin County, Minnesota"	"3301"	"96875"
[6,]	"Census Tract 11, Hennepin County, Minnesota"	"2004"	"69509"

	state	county	tract
[1,]	"27"	"053"	"000101"
[2,]	"27"	"053"	"000102"

```
[3,] "27" "053" "000300"
[4,] "27" "053" "000601"
[5,] "27" "053" "000603"
[6,] "27" "053" "001100"
```

```
hennepin_htr2 <- as_tibble(resp_json_df) |>
  mutate(population = parse_number(B01003_001E),
         median_income = parse_number(B19013_001E)) |>
  select(-B01003_001E, -B19013_001E, -state, -county)

hennepin_htr2 |>
  ggplot(aes(x = population, y = median_income)) +
    geom_point() +
    geom_smooth(method = "lm")
```



```
summary(hennepin_htr2$population)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0	2876	3714	3815	4651	9680

```
summary(hennepin_httr2$median_income)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-6666666666	61354	80966	-3966166	107232	250001

```
sort(hennepin_httr2$population)
```

```
[1] 0 223 1514 1622 1672 1760 1766 1779 1798 1844 1848 1877 1897 1915 1926
[16] 1935 1942 1973 2000 2004 2012 2013 2017 2038 2058 2061 2067 2092 2111 2123
[31] 2130 2150 2163 2228 2235 2256 2272 2274 2280 2283 2295 2315 2339 2341 2357
[46] 2399 2415 2416 2419 2460 2462 2476 2484 2499 2511 2511 2528 2532 2551 2570
[61] 2594 2605 2625 2656 2658 2668 2670 2675 2681 2724 2738 2756 2763 2780 2796
[76] 2808 2820 2822 2837 2848 2853 2865 2876 2878 2916 2935 2944 2950 2954 2969
[91] 2971 2984 2994 3001 3036 3037 3038 3046 3047 3048 3075 3077 3119 3124 3127
[106] 3138 3150 3152 3162 3168 3193 3222 3224 3224 3225 3236 3251 3274 3298 3301
[121] 3305 3317 3317 3326 3331 3335 3341 3364 3372 3376 3379 3386 3404 3404 3418
[136] 3431 3439 3444 3454 3466 3472 3474 3486 3498 3512 3513 3557 3573 3574 3575
[151] 3585 3607 3628 3631 3634 3654 3656 3666 3671 3673 3676 3687 3703 3710 3714
[166] 3739 3750 3762 3764 3765 3799 3801 3805 3806 3808 3810 3811 3829 3832 3842
[181] 3853 3862 3877 3885 3890 3895 3896 3896 3903 3903 3913 3924 3930 3960 3967
[196] 3972 3974 3976 3978 3980 3989 3995 4008 4010 4013 4025 4036 4063 4086 4097
[211] 4098 4126 4132 4179 4200 4219 4228 4237 4273 4286 4295 4305 4319 4321 4326
[226] 4355 4359 4366 4371 4378 4385 4412 4441 4455 4460 4466 4472 4481 4503 4535
[241] 4584 4587 4591 4613 4622 4629 4651 4665 4671 4678 4693 4696 4706 4713 4718
[256] 4728 4747 4767 4769 4789 4789 4815 4855 4855 4874 4899 4919 4930 4972 4978
[271] 4983 4992 5030 5033 5041 5065 5085 5099 5107 5150 5195 5213 5244 5262 5267
[286] 5295 5305 5313 5364 5366 5385 5386 5415 5442 5459 5507 5510 5515 5541 5541
[301] 5587 5709 5725 5781 5821 5831 5872 5880 5980 6025 6069 6071 6102 6113 6166
[316] 6229 6249 6258 6265 6308 6482 6595 6709 6928 7286 7604 7828 9486 9680
```

```
sort(hennepin_httr2$median_income)
```

[1]	-6666666666	-6666666666	14748	20000	22768	23256
[7]	23391	25708	31513	31981	32321	32758
[13]	34273	35368	35855	36700	37315	37346
[19]	37413	38286	38554	39420	39605	39609
[25]	39630	40400	40476	40603	40867	42426
[31]	42550	42753	43036	43750	44867	45640
[37]	46157	46333	46596	47139	47197	47688
[43]	47857	48464	48690	48750	49028	49139

[49]	49659	50000	50741	50755	50935	51250
[55]	51513	51705	51923	52169	52304	52370
[61]	52781	52917	53393	53542	53564	53952
[67]	54026	54636	55321	55430	55833	56338
[73]	56955	57469	57802	57875	58426	59013
[79]	59704	59876	60375	61213	61354	61547
[85]	62188	62279	62404	62426	62770	63750
[91]	63990	64250	64333	64621	64676	64792
[97]	65323	65329	65395	65455	65590	65772
[103]	66364	66452	66549	66875	67102	67132
[109]	67473	67614	68114	68158	68369	68417
[115]	68434	68796	68913	68971	69509	69600
[121]	70089	70927	70970	71071	71122	71146
[127]	71250	71670	71818	72054	72102	72766
[133]	72853	73482	73514	73527	73897	73984
[139]	74286	74330	74817	75147	75556	75833
[145]	76111	76164	76417	76792	76839	77500
[151]	78137	78171	78333	78418	78509	78605
[157]	78728	79167	79191	79366	79750	80012
[163]	80080	80350	80966	81341	81341	81411
[169]	81977	82014	82098	82340	82527	83090
[175]	83250	83315	83380	84063	84569	84583
[181]	84792	85078	85221	85938	86106	86111
[187]	86904	87054	87390	87426	87599	87857
[193]	88431	88542	88895	89417	89740	89792
[199]	89891	89922	90167	91230	91250	91333
[205]	91637	91827	92019	92683	92941	93011
[211]	93750	94656	95750	95855	95980	96328
[217]	96378	96667	96856	96875	96983	97609
[223]	98137	98550	98986	99792	99853	100054
[229]	100329	100652	100761	101156	101194	101440
[235]	101578	103049	103531	103611	103750	104242
[241]	104306	104412	104795	104904	106310	106518
[247]	107232	107303	108476	108510	109722	110125
[253]	110339	110694	110729	110774	111364	111635
[259]	111950	112104	112557	112566	113563	113750
[265]	114550	115934	116281	116861	117631	118333
[271]	118594	118697	118828	119214	119821	120769
[277]	122180	122206	123312	125750	126250	127375
[283]	127396	130404	130486	131023	131042	132361
[289]	132604	133333	133472	133504	133859	134250
[295]	136012	136369	138848	141528	141984	142500
[301]	142889	143125	143744	143935	144282	144318

[307]	146328	147237	147672	148512	148611	149934
[313]	153917	154306	159857	161458	161471	165865
[319]	176580	178259	179743	179926	180463	185357
[325]	194417	194882	200438	202098	250001	

```


hennepin_htr2 <- hennepin_htr2 |>
  mutate(median_income = ifelse(median_income > 0, median_income, NA),
         population = ifelse(population > 0, population, NA))

hennepin_htr2 |>
  ggplot(aes(x = population, y = median_income)) +
    geom_point() +
    geom_smooth(method = "lm")

```

Warning: Removed 2 rows containing non-finite outside the scale range (`stat_smooth()`).

Warning: Removed 2 rows containing missing values or values outside the scale range (`geom_point()`).



```
# To make choropleth map by census tract, would need to download US Census
# Bureau TIGER geometries using tigris package
```

```
# Finally using httr
url <- str_c("https://api.census.gov/data/2020/acs/acs5?get=NAME,B01003_001E,B19013_001E&for=")
acs5 <- GET(url)
details <- content(acs5, "parsed")
# details
details[[1]] # variable names
```

```
[[1]]
[1] "NAME"

[[2]]
[1] "B01003_001E"

[[3]]
[1] "B19013_001E"

[[4]]
[1] "state"

[[5]]
[1] "county"

[[6]]
[1] "tract"
```

```
details[[2]] # list with information on 1st tract
```

```
[[1]]
[1] "Census Tract 1.01, Hennepin County, Minnesota"

[[2]]
[1] "3472"

[[3]]
[1] "70927"

[[4]]
[1] "27"
```



```
[[5]]  
[1] "053"
```

```
[[6]]  
[1] "000101"
```

```
name = character()  
population = double()  
median_income = double()  
tract = character()  
  
for(i in 2:330) {  
  name[i-1] <- details[[i]][[1]][1]  
  population[i-1] <- details[[i]][[2]][1]  
  median_income[i-1] <- details[[i]][[3]][1]  
  tract[i-1] <- details[[i]][[6]][1]  
}  
hennepin_httr <- tibble(  
  name = name,  
  population = parse_number(population),  
  median_income = parse_number(median_income),  
  tract = tract  
)
```

On Your Own

1. Write a for loop to obtain the Hennepin County data from 2017-2021
2. Write a function to give choices about year, county, and variables

```
# function to allow user inputs  
  
MN_tract_data <- function(year, county, variables) {  
  tidycensus::get_acs(  
    Sys.sleep(0.5),  
    year = year,  
    state = "MN",  
    geography = "tract",  
    variables = variables,  
    output = "wide",  
    geometry = TRUE,  
  )  
}
```

```

    county = county
  ) |>
    mutate(year = year)
}

# Should really build in checks so that county is in MN, year is in
#   proper range, and variables are part of ACS1 data set

my_data <- MN_tract_data(year = 2021,
  county = "Hennepin",
  variables = c("B01003_001", "B19013_001"))

```


Getting data from the 2017-2021 5-year ACS

Downloading feature geometry from the Census website. To cache shapefiles for use in future

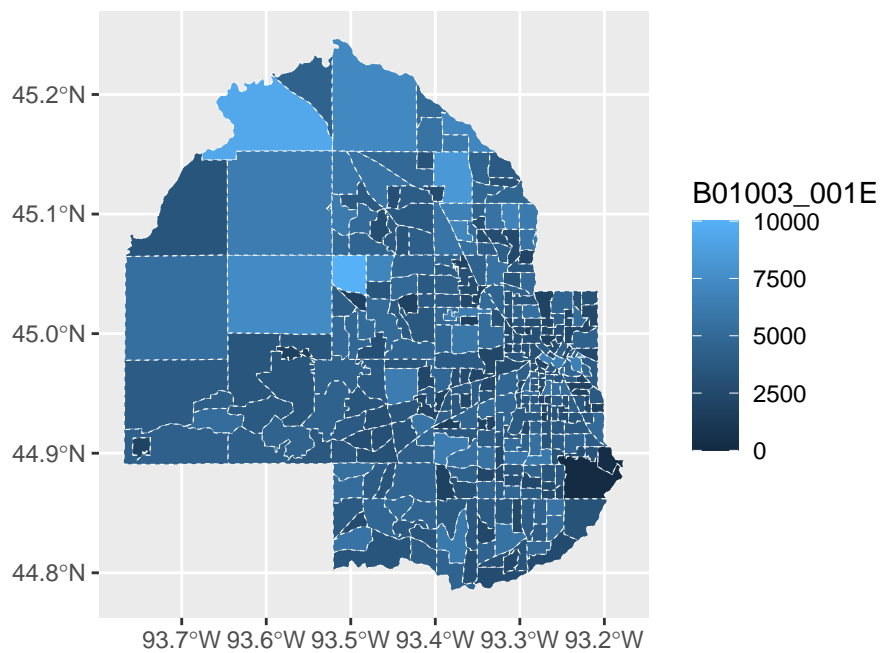
		0%
=		2%
=====		7%
=====		9%
=====		14%
=====		16%
=====		17%
=====		18%
=====		21%
=====		22%
=====		24%
=====		25%

=====		27%
=====		28%
=====		29%
=====		31%
=====		32%
=====		33%
=====		35%
=====		36%
=====		38%
=====		40%
=====		41%
=====		42%
=====		44%
=====		46%
=====		48%
=====		49%
=====		51%
=====		52%
=====		55%
=====		56%
=====		58%

=====	59%
=====	60%
=====	62%
=====	63%
=====	64%
=====	65%
=====	67%
=====	68%
=====	69%
=====	71%
=====	72%
=====	73%
=====	75%
=====	75%
=====	76%
=====	78%
=====	80%
=====	82%
=====	83%
=====	84%
=====	86%
=====	87%



```
ggplot(data = my_data) +
  geom_sf(aes(fill = B01003_001E), colour = "white", linetype = 2)
```



```
my_data <- MN_tract_data(year = 2022,
  county = "Rice",
  variables = c("B01003_001", "B19013_001"))
```

Getting data from the 2018-2022 5-year ACS

Downloading feature geometry from the Census website. To cache shapefiles for use in future


	0%
=	2%
==	3%
====	5%
=====	7%
=====	9%
=====	13%
=====	14%
=====	16%
=====	17%
=====	18%
=====	21%
=====	24%
=====	26%
=====	27%
=====	29%
=====	30%
=====	31%
=====	34%

=====		35%
=====		37%
=====		38%
=====		39%
=====		40%
=====		42%
=====		44%
=====		46%
=====		47%
=====		48%
=====		50%
=====		51%
=====		52%
=====		53%
=====		55%
=====		56%
=====		57%
=====		59%
=====		60%
=====		61%
=====		63%

=====	64%
=====	65%
=====	66%
=====	68%
=====	69%
=====	70%
=====	72%
=====	73%
=====	74%
=====	76%
=====	77%
=====	78%
=====	79%
=====	81%
=====	82%
=====	83%
=====	85%
=====	86%
=====	87%
=====	90%
=====	92%
=====	95%


```
|
|=====| 96%
|
|=====| 98%
|
|=====| 99%
|
|=====| 100%
```

```
ggplot(data = my_data) +
  geom_sf(aes(fill = B01003_001E), colour = "white", linetype = 2)
```



```
# Try other variables:
# - B25077_001 is median home price
# - B02001_002 is number of white residents
# - etc.
# although the census codebook is admittedly quite daunting!
```

3. Use your function from (2) along with `map` and `list_rbind` to build a data set for Rice county for the years 2019-2021

```
# To examine trends over time in Rice County
2019:2021 |>
  purrr::map(\(x)
    MN_tract_data(
      x,
      county = "Rice",
      variables = c("B01003_001", "B19013_001")
    )
  ) |>
  list_rbind()
```

Getting data from the 2015-2019 5-year ACS

Downloading feature geometry from the Census website. To cache shapefiles for use in future

	0%
=	2%
===	4%
====	6%
=====	8%
=====	16%
=====	18%
=====	19%
=====	20%
=====	22%
=====	23%
=====	25%

=====	29%
=====	30%
=====	32%
=====	33%
=====	35%
=====	37%
=====	39%
=====	41%
=====	43%
=====	45%
=====	47%
=====	49%
=====	51%
=====	52%
=====	54%
=====	57%
=====	59%
=====	62%
=====	63%
=====	66%
=====	68%
=====	70%



Getting data from the 2016-2020 5-year ACS

Downloading feature geometry from the Census website. To cache shapefiles for use in future

Getting data from the 2017-2021 5-year ACS

Downloading feature geometry from the Census website. To cache shapefiles for use in future

GEOID

NAME B01003_001E

1	27131070504	Census Tract 705.04, Rice County, Minnesota	3933
2	27131070400	Census Tract 704, Rice County, Minnesota	4511
3	27131070300	Census Tract 703, Rice County, Minnesota	4551
4	27131070503	Census Tract 705.03, Rice County, Minnesota	3348
5	27131070601	Census Tract 706.01, Rice County, Minnesota	3526
6	27131070800	Census Tract 708, Rice County, Minnesota	8101
7	27131070901	Census Tract 709.01, Rice County, Minnesota	5509
8	27131070700	Census Tract 707, Rice County, Minnesota	7165
9	27131070100	Census Tract 701, Rice County, Minnesota	7333
10	27131070602	Census Tract 706.02, Rice County, Minnesota	5211
11	27131070200	Census Tract 702, Rice County, Minnesota	5463
12	27131070902	Census Tract 709.02, Rice County, Minnesota	3160
13	27131070501	Census Tract 705.01, Rice County, Minnesota	4374
14	27131070501	Census Tract 705.01, Rice County, Minnesota	4272
15	27131070504	Census Tract 705.04, Rice County, Minnesota	3941
16	27131070801	Census Tract 708.01, Rice County, Minnesota	4456
17	27131070200	Census Tract 702, Rice County, Minnesota	5508
18	27131070701	Census Tract 707.01, Rice County, Minnesota	3057
19	27131070400	Census Tract 704, Rice County, Minnesota	4686
20	27131070300	Census Tract 703, Rice County, Minnesota	4737
21	27131070601	Census Tract 706.01, Rice County, Minnesota	3669
22	27131070102	Census Tract 701.02, Rice County, Minnesota	3786
23	27131070802	Census Tract 708.02, Rice County, Minnesota	3873
24	27131070702	Census Tract 707.02, Rice County, Minnesota	3872
25	27131070901	Census Tract 709.01, Rice County, Minnesota	5681
26	27131070503	Census Tract 705.03, Rice County, Minnesota	3185
27	27131070902	Census Tract 709.02, Rice County, Minnesota	2992
28	27131070101	Census Tract 701.01, Rice County, Minnesota	3428
29	27131070602	Census Tract 706.02, Rice County, Minnesota	5406
30	27131070902	Census Tract 709.02, Rice County, Minnesota	3212
31	27131070601	Census Tract 706.01, Rice County, Minnesota	3775
32	27131070503	Census Tract 705.03, Rice County, Minnesota	3035
33	27131070702	Census Tract 707.02, Rice County, Minnesota	3738
34	27131070901	Census Tract 709.01, Rice County, Minnesota	5858
35	27131070801	Census Tract 708.01, Rice County, Minnesota	4618
36	27131070501	Census Tract 705.01, Rice County, Minnesota	4242
37	27131070300	Census Tract 703, Rice County, Minnesota	4657
38	27131070200	Census Tract 702, Rice County, Minnesota	5419
39	27131070400	Census Tract 704, Rice County, Minnesota	4380
40	27131070701	Census Tract 707.01, Rice County, Minnesota	3028
41	27131070504	Census Tract 705.04, Rice County, Minnesota	3917
42	27131070101	Census Tract 701.01, Rice County, Minnesota	3417
43	27131070802	Census Tract 708.02, Rice County, Minnesota	3944

44	27131070102	Census Tract 701.02, Rice County, Minnesota	4201
45	27131070602	Census Tract 706.02, Rice County, Minnesota	5354
	B01003_001M	B19013_001E B19013_001M	geometry year
1	273	63989	9273 MULTIPOLYGON (((-93.19137 4... 2019
2	168	85952	2758 MULTIPOLYGON (((-93.40564 4... 2019
3	190	78343	4242 MULTIPOLYGON (((-93.52521 4... 2019
4	245	92321	14200 MULTIPOLYGON (((-93.16075 4... 2019
5	333	50368	9979 MULTIPOLYGON (((-93.17615 4... 2019
6	465	48403	7679 MULTIPOLYGON (((-93.29819 4... 2019
7	456	44417	10552 MULTIPOLYGON (((-93.30904 4... 2019
8	414	67868	9422 MULTIPOLYGON (((-93.27265 4... 2019
9	326	91667	8106 MULTIPOLYGON (((-93.52452 4... 2019
10	310	64479	12376 MULTIPOLYGON (((-93.22644 4... 2019
11	177	101359	4104 MULTIPOLYGON (((-93.5246 44... 2019
12	410	45230	12887 MULTIPOLYGON (((-93.30888 4... 2019
13	270	66188	9179 MULTIPOLYGON (((-93.16981 4... 2019
14	316	64792	13256 MULTIPOLYGON (((-93.16981 4... 2020
15	536	63500	7351 MULTIPOLYGON (((-93.1909 44... 2020
16	703	67625	23325 MULTIPOLYGON (((-93.29829 4... 2020
17	473	104011	5648 MULTIPOLYGON (((-93.5246 44... 2020
18	218	73750	13139 MULTIPOLYGON (((-93.26704 4... 2020
19	296	86094	3438 MULTIPOLYGON (((-93.40564 4... 2020
20	244	79068	4902 MULTIPOLYGON (((-93.52518 4... 2020
21	525	52936	10436 MULTIPOLYGON (((-93.17615 4... 2020
22	199	96023	13649 MULTIPOLYGON (((-93.44292 4... 2020
23	437	63924	8715 MULTIPOLYGON (((-93.28272 4... 2020
24	425	49811	16864 MULTIPOLYGON (((-93.27265 4... 2020
25	566	51595	9615 MULTIPOLYGON (((-93.30904 4... 2020
26	341	100516	11630 MULTIPOLYGON (((-93.16075 4... 2020
27	440	46750	15457 MULTIPOLYGON (((-93.30888 4... 2020
28	295	100563	15809 MULTIPOLYGON (((-93.52452 4... 2020
29	377	62078	5270 MULTIPOLYGON (((-93.22644 4... 2020
30	421	47059	15456 MULTIPOLYGON (((-93.30888 4... 2021
31	435	56319	4333 MULTIPOLYGON (((-93.17615 4... 2021
32	321	105952	8429 MULTIPOLYGON (((-93.16075 4... 2021
33	409	57126	13968 MULTIPOLYGON (((-93.27265 4... 2021
34	714	47344	9579 MULTIPOLYGON (((-93.30904 4... 2021
35	622	61193	23977 MULTIPOLYGON (((-93.29829 4... 2021
36	380	79063	15272 MULTIPOLYGON (((-93.16981 4... 2021
37	296	83911	7244 MULTIPOLYGON (((-93.52522 4... 2021
38	520	111711	10313 MULTIPOLYGON (((-93.5246 44... 2021
39	274	90179	4919 MULTIPOLYGON (((-93.40564 4... 2021
40	358	82500	20934 MULTIPOLYGON (((-93.26775 4... 2021

```

41          537          67219          9805 MULTIPOLYGON (((-93.1909 44... 2021
42          270          108490          1768 MULTIPOLYGON (((-93.52452 4... 2021
43          462          63679          12261 MULTIPOLYGON (((-93.28274 4... 2021
44          199          85789          20094 MULTIPOLYGON (((-93.44292 4... 2021
45          359          63835          4805 MULTIPOLYGON (((-93.22644 4... 2021

```

```

# Or a little more simply
2019:2021 |>
  purrr::map(MN_tract_data,
    county = "Rice",
    variables = c("B01003_001", "B19013_001")
  ) |>
  list_rbind()

```

Getting data from the 2015-2019 5-year ACS

Downloading feature geometry from the Census website. To cache shapefiles for use in future

Getting data from the 2016-2020 5-year ACS

Downloading feature geometry from the Census website. To cache shapefiles for use in future

Getting data from the 2017-2021 5-year ACS

Downloading feature geometry from the Census website. To cache shapefiles for use in future

	GEOID	NAME	B01003_001E
1	27131070504	Census Tract 705.04, Rice County, Minnesota	3933
2	27131070400	Census Tract 704, Rice County, Minnesota	4511
3	27131070300	Census Tract 703, Rice County, Minnesota	4551
4	27131070503	Census Tract 705.03, Rice County, Minnesota	3348
5	27131070601	Census Tract 706.01, Rice County, Minnesota	3526
6	27131070800	Census Tract 708, Rice County, Minnesota	8101
7	27131070901	Census Tract 709.01, Rice County, Minnesota	5509
8	27131070700	Census Tract 707, Rice County, Minnesota	7165
9	27131070100	Census Tract 701, Rice County, Minnesota	7333
10	27131070602	Census Tract 706.02, Rice County, Minnesota	5211
11	27131070200	Census Tract 702, Rice County, Minnesota	5463
12	27131070902	Census Tract 709.02, Rice County, Minnesota	3160
13	27131070501	Census Tract 705.01, Rice County, Minnesota	4374
14	27131070501	Census Tract 705.01, Rice County, Minnesota	4272

15	27131070504	Census Tract 705.04, Rice County, Minnesota	3941
16	27131070801	Census Tract 708.01, Rice County, Minnesota	4456
17	27131070200	Census Tract 702, Rice County, Minnesota	5508
18	27131070701	Census Tract 707.01, Rice County, Minnesota	3057
19	27131070400	Census Tract 704, Rice County, Minnesota	4686
20	27131070300	Census Tract 703, Rice County, Minnesota	4737
21	27131070601	Census Tract 706.01, Rice County, Minnesota	3669
22	27131070102	Census Tract 701.02, Rice County, Minnesota	3786
23	27131070802	Census Tract 708.02, Rice County, Minnesota	3873
24	27131070702	Census Tract 707.02, Rice County, Minnesota	3872
25	27131070901	Census Tract 709.01, Rice County, Minnesota	5681
26	27131070503	Census Tract 705.03, Rice County, Minnesota	3185
27	27131070902	Census Tract 709.02, Rice County, Minnesota	2992
28	27131070101	Census Tract 701.01, Rice County, Minnesota	3428
29	27131070602	Census Tract 706.02, Rice County, Minnesota	5406
30	27131070902	Census Tract 709.02, Rice County, Minnesota	3212
31	27131070601	Census Tract 706.01, Rice County, Minnesota	3775
32	27131070503	Census Tract 705.03, Rice County, Minnesota	3035
33	27131070702	Census Tract 707.02, Rice County, Minnesota	3738
34	27131070901	Census Tract 709.01, Rice County, Minnesota	5858
35	27131070801	Census Tract 708.01, Rice County, Minnesota	4618
36	27131070501	Census Tract 705.01, Rice County, Minnesota	4242
37	27131070300	Census Tract 703, Rice County, Minnesota	4657
38	27131070200	Census Tract 702, Rice County, Minnesota	5419
39	27131070400	Census Tract 704, Rice County, Minnesota	4380
40	27131070701	Census Tract 707.01, Rice County, Minnesota	3028
41	27131070504	Census Tract 705.04, Rice County, Minnesota	3917
42	27131070101	Census Tract 701.01, Rice County, Minnesota	3417
43	27131070802	Census Tract 708.02, Rice County, Minnesota	3944
44	27131070102	Census Tract 701.02, Rice County, Minnesota	4201
45	27131070602	Census Tract 706.02, Rice County, Minnesota	5354
	B01003_001M	B19013_001E B19013_001M	geometry year
1	273	63989	9273 MULTIPOLYGON (((-93.19137 4... 2019
2	168	85952	2758 MULTIPOLYGON (((-93.40564 4... 2019
3	190	78343	4242 MULTIPOLYGON (((-93.52521 4... 2019
4	245	92321	14200 MULTIPOLYGON (((-93.16075 4... 2019
5	333	50368	9979 MULTIPOLYGON (((-93.17615 4... 2019
6	465	48403	7679 MULTIPOLYGON (((-93.29819 4... 2019
7	456	44417	10552 MULTIPOLYGON (((-93.30904 4... 2019
8	414	67868	9422 MULTIPOLYGON (((-93.27265 4... 2019
9	326	91667	8106 MULTIPOLYGON (((-93.52452 4... 2019
10	310	64479	12376 MULTIPOLYGON (((-93.22644 4... 2019
11	177	101359	4104 MULTIPOLYGON (((-93.5246 44... 2019

12	410	45230	12887	MULTIPOLYGON	(((−93.30888 4... 2019
13	270	66188	9179	MULTIPOLYGON	(((−93.16981 4... 2019
14	316	64792	13256	MULTIPOLYGON	(((−93.16981 4... 2020
15	536	63500	7351	MULTIPOLYGON	(((−93.1909 44... 2020
16	703	67625	23325	MULTIPOLYGON	(((−93.29829 4... 2020
17	473	104011	5648	MULTIPOLYGON	(((−93.5246 44... 2020
18	218	73750	13139	MULTIPOLYGON	(((−93.26704 4... 2020
19	296	86094	3438	MULTIPOLYGON	(((−93.40564 4... 2020
20	244	79068	4902	MULTIPOLYGON	(((−93.52518 4... 2020
21	525	52936	10436	MULTIPOLYGON	(((−93.17615 4... 2020
22	199	96023	13649	MULTIPOLYGON	(((−93.44292 4... 2020
23	437	63924	8715	MULTIPOLYGON	(((−93.28272 4... 2020
24	425	49811	16864	MULTIPOLYGON	(((−93.27265 4... 2020
25	566	51595	9615	MULTIPOLYGON	(((−93.30904 4... 2020
26	341	100516	11630	MULTIPOLYGON	(((−93.16075 4... 2020
27	440	46750	15457	MULTIPOLYGON	(((−93.30888 4... 2020
28	295	100563	15809	MULTIPOLYGON	(((−93.52452 4... 2020
29	377	62078	5270	MULTIPOLYGON	(((−93.22644 4... 2020
30	421	47059	15456	MULTIPOLYGON	(((−93.30888 4... 2021
31	435	56319	4333	MULTIPOLYGON	(((−93.17615 4... 2021
32	321	105952	8429	MULTIPOLYGON	(((−93.16075 4... 2021
33	409	57126	13968	MULTIPOLYGON	(((−93.27265 4... 2021
34	714	47344	9579	MULTIPOLYGON	(((−93.30904 4... 2021
35	622	61193	23977	MULTIPOLYGON	(((−93.29829 4... 2021
36	380	79063	15272	MULTIPOLYGON	(((−93.16981 4... 2021
37	296	83911	7244	MULTIPOLYGON	(((−93.52522 4... 2021
38	520	111711	10313	MULTIPOLYGON	(((−93.5246 44... 2021
39	274	90179	4919	MULTIPOLYGON	(((−93.40564 4... 2021
40	358	82500	20934	MULTIPOLYGON	(((−93.26775 4... 2021
41	537	67219	9805	MULTIPOLYGON	(((−93.1909 44... 2021
42	270	108490	1768	MULTIPOLYGON	(((−93.52452 4... 2021
43	462	63679	12261	MULTIPOLYGON	(((−93.28274 4... 2021
44	199	85789	20094	MULTIPOLYGON	(((−93.44292 4... 2021
45	359	63835	4805	MULTIPOLYGON	(((−93.22644 4... 2021

One more example using an API key

Here's an example of getting data from a website that attempts to make imdb movie data available as an API.

Initial instructions:

- go to omdbapi.com under the API Key tab and request a free API key

- store your key as discussed earlier
- explore the examples at omdbapi.com

We will first obtain data about the movie Coco from 2017.

```
# I added: Sys.setenv(OMDB_KEY = "")
# I used the first line to store my OMDB API key in .Renviron
# Sys.setenv(OMDB_KEY = "paste my omdb key here")
myapikey <- Sys.getenv("OMDB_KEY")

# Find url exploring examples at omdbapi.com
url <- str_c("http://www.omdbapi.com/?t=Coco&y=2017&apikey=", myapikey)

coco <- GET(url)    # coco holds response from server
coco              # Status of 200 is good!
```

```
Response [http://www.omdbapi.com/?t=Coco&y=2017&apikey=d64645f3]
  Date: 2025-03-26 00:23
  Status: 200
  Content-Type: application/json; charset=utf-8
  Size: 1.04 kB
```

```
details <- content(coco, "parse")
details                                     # get a list of 25 pieces of information
```

```
$Title
[1] "Coco"
```

```
$Year
[1] "2017"
```

```
$Rated
[1] "PG"
```

```
$Released
[1] "22 Nov 2017"
```

```
$Runtime
[1] "105 min"
```

```
$Genre
```

[1] "Animation, Adventure, Drama"

\$Director

[1] "Lee Unkrich, Adrian Molina"

\$Writer

[1] "Lee Unkrich, Jason Katz, Matthew Aldrich"

\$Actors

[1] "Anthony Gonzalez, Gael García Bernal, Benjamin Bratt"

\$Plot

[1] "Aspiring musician Miguel, confronted with his family's ancestral ban on music, enters tl

\$Language

[1] "English, Spanish"

\$Country

[1] "United States, Mexico"

\$Awards

[1] "Won 2 Oscars. 112 wins & 42 nominations total"

\$Poster

[1] "https://m.media-amazon.com/images/M/MV5BMDIyM2E2NTAtMzlhNy00ZGUxLWI1NjgtZDY5MzhiMDc5NGU"

\$Ratings

\$Ratings[[1]]

\$Ratings[[1]]\$Source

[1] "Internet Movie Database"

\$Ratings[[1]]\$Value

[1] "8.4/10"

\$Ratings[[2]]

\$Ratings[[2]]\$Source

[1] "Rotten Tomatoes"

\$Ratings[[2]]\$Value

[1] "97%"

```
$Ratings[[3]]  
$Ratings[[3]]$Source  
[1] "Metacritic"
```

```
$Ratings[[3]]$Value  
[1] "81/100"
```

```
$Metascore  
[1] "81"
```

```
$imdbRating  
[1] "8.4"
```

```
$imdbVotes  
[1] "635,840"
```

```
$imdbID  
[1] "tt2380307"
```

```
$Type  
[1] "movie"
```

```
$DVD  
[1] "N/A"
```

```
$BoxOffice  
[1] "$210,460,015"
```

```
$Production  
[1] "N/A"
```

```
$Website  
[1] "N/A"
```

```
$Response  
[1] "True"
```

```
details$Year # how to access details
```

```
[1] "2017"
```

```
details[[2]] # since a list, another way to access
```

```
[1] "2017"
```

Now build a data set for a collection of movies

```
# Must figure out pattern in URL for obtaining different movies
# - try searching for others
movies <- c("Coco", "Wonder+Woman", "Get+Out",
           "The+Greatest+Showman", "Thor:+Ragnarok")

# Set up empty tibble
omdb <- tibble(Title = character(), Rated = character(), Genre = character(),
              Actors = character(), Metascore = double(), imdbRating = double(),
              BoxOffice = double())

# Use for loop to run through API request process 5 times,
# each time filling the next row in the tibble
# - can do max of 1000 GETs per day
for(i in 1:5) {
  url <- str_c("http://www.omdbapi.com/?t=", movies[i],
              "&apikey=", myapikey)
  Sys.sleep(0.5)
  onemovie <- GET(url)
  details <- content(onemovie, "parse")
  omdb[i,1] <- details$Title
  omdb[i,2] <- details$Rated
  omdb[i,3] <- details$Genre
  omdb[i,4] <- details$Actors
  omdb[i,5] <- parse_number(details$Metascore)
  omdb[i,6] <- parse_number(details$imdbRating)
  omdb[i,7] <- parse_number(details$BoxOffice) # no $ and , 's
}

omdb
```

```
# A tibble: 5 x 7
```

	Title <chr>	Rated <chr>	Genre <chr>	Actors <chr>	Metascore <dbl>	imdbRating <dbl>	BoxOffice <dbl>
1	Coco	PG	Animation, A~	Antho~	81	8.4	210460015
2	Wonder Woman	PG-13	Action, Adve~	Gal G~	76	7.3	412845172

3	Get Out	R	Horror, Myst~ Danie~	85	7.8	176196665
4	The Greatest Showman	PG	Biography, D~ Hugh ~	48	7.5	174340174
5	Thor: Ragnarok	PG-13	Action, Adve~ Chris~	74	7.9	315058289

```
# could use stringr functions to further organize this data - separate
# different genres, different actors, etc.
```

```
movies <- c("Up", "Cars", "Kung+Fu+Panda", "The+Emperor%27s+New+Groove", "Mulan")

omdb <- tibble(Title = character(), Released = character(), Runtime = character(), Plot = cha

for(i in 1:5) {
  url <- str_c("http://www.omdbapi.com/?t=",movies[i],
              "&apikey=", myapikey)
  Sys.sleep(0.5)
  onemovie <- GET(url)
  details <- content(onemovie, "parse")
  omdb[i,1] <- details$Title
  omdb[i,2] <- details$Released
  omdb[i,3] <- details$Runtime
  omdb[i,4] <- details$Plot
  omdb[i,5] <- parse_number(details$BoxOffice)
}
```

On Your Own (continued)

4. (Based on final project by Mary Wu and Jenna Graff, MSCS 264, Spring 2024). Start with a small data set on 56 national parks from [kaggle](#), and supplement with columns for the park address (a single column including address, city, state, and zip code) and a list of available activities (a single character column with activities separated by commas) from the park websites themselves.

Preliminaries:

- Request API [here](#)
- Check out [API guide](#)

```
np_kaggle <- read_csv("Data/parks.csv")
```

You can download this .qmd file from [here](#). Just hit the Download Raw File button.

Using rvest for web scraping

If you would like to assemble data from a website with no API, you can often acquire data using more brute force methods commonly called web scraping. Typically, this involves finding content inside HTML (Hypertext markup language) code used for creating webpages and web applications and the CSS (Cascading style sheets) language for customizing the appearance of webpages. We are used to reading data from .csv files.... but most websites have it stored in XML (like html, but for data). You can read more about it here if you're interested: <https://www.w3schools.com/xml/default.asp>

XML has a sort of tree or graph-like structure... so we can identify information by which `node` it belongs to (`html_nodes`) and then convert the content into something we can use in R (`html_text` or `html_table`).

Here's one quick example of web scraping. First check out the webpage https://www.cheese.com/by_type and then select Semi-Soft. We can drill into the html code for this webpage and find and store specific information (like cheese names)

```
session <- bow("https://www.cheese.com/by_type", force = TRUE)
result <- scrape(session, query=list(t="semi-soft", per_page=100)) |>
  html_node("#main-body") |>
  html_nodes("h3") |>
  html_text()
head(result)
```

```
[1] "American Cheese"      "Mozzarella"          "Taleggio"
[4] "Fontina Val d'Aosta"  "Blue Cheese"         "Jarlsberg"
```

```
#> [1] "3-Cheese Italian Blend" "Abbaye de Citeaux"
#> [3] "Abbaye du Mont des Cats" "Adelost"
#> [5] "ADL Brick Cheese"      "Ailsa Craig"
```

Four steps to scraping data with functions in the rvest library:

0. `robotstxt::paths_allowed()` Check if the website allows scraping, and then make sure we scrape “politely”
1. `read_html()`. Input the URL containing the data and turn the html code into an XML file (another markup format that's easier to work with).
2. `html_nodes()`. Extract specific nodes from the XML file by using the CSS path that leads to the content of interest. (use `css="table"` for tables.)
3. `html_text()`. Extract content of interest from nodes. Might also use `html_table()` etc.

Data scraping ethics

Before scraping, we should always check first whether the website allows scraping. We should also consider if there's any personal or confidential information, and we should be considerate to not overload the server we're scraping from.

[Chapter 24 in R4DS](#) provides a nice overview of some of the important issues to consider. A couple of highlights:

- be aware of terms of service, and, if available, the `robots.txt` file that some websites will publish to clarify what can and cannot be scraped and other constraints about scraping.
- use the [polite package](#) to scrape public, non-personal, and factual data in a respectful manner
- scrape with a good purpose and request only what you need; in particular, be extremely wary of personally identifiable information

See [this article](#) for more perspective on the ethics of data scraping.

When the data is already in table form:

In this example, we will scrape climate data from [this website](#)

The website already contains data in table form, so we use `html_nodes(. , css = "table")` and `html_table()`

```
# check that scraping is allowed (Step 0)
robotstxt::paths_allowed("https://www.usclimatedata.com/climate/minneapolis/minnesota/united-states")
```

```
www.usclimatedata.com
```

```
[1] TRUE
```

```
# Step 1: read_html()
mpls <- read_html("https://www.usclimatedata.com/climate/minneapolis/minnesota/united-states")

# 2: html_nodes()
tables <- html_nodes(mpls, css = "table")
tables # have to guesstimate which table contains climate info
```



```
{xml_nodeset (8)}
[1] <table id="monthly_table_one" class="table table-hover tablesaw tablesaw- ...
[2] <table id="monthly_table_two" class="table table-hover tablesaw tablesaw- ...
[3] <table class="table table-hover tablesaw tablesaw-mode-swipe mt-4 daily_t ...
[4] <table class="table table-hover tablesaw tablesaw-mode-swipe mt-4 history ...
[5] <table class="table table-striped table-hover tablesaw tablesaw-mode-swip ...
[6] <table class="table table-hover tablesaw geo_table">\n<thead><tr>\n<th> < ...
[7] <table class="table table-hover tablesaw datetime_table" data-tablesaw-hi ...
[8] <table class="table table-hover tablesaw monthly_summary_table" data-tabl ...
```

```
# 3: html_table()
html_table(tables, header = TRUE, fill = TRUE) # find the right table
```

```
[[1]]
# A tibble: 6 x 7
  <chr> JanJa FebFe MarMa AprAp MayMa JunJu
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 Average high in °F Av. high Hi 24 29 41 58 69 79
2 Average low in °F Av. low Lo 8 13 24 37 49 59
3 Days with precipitation Days precip.~ 8 7 11 9 11 13
4 Hours of sunshine Hours sun. Sun 140 166 200 231 272 302
5 Av. precipitation in inch Av. precip~ 0.9 0.77 1.89 2.66 3.36 4.25
6 Av. snowfall in inch Snowfall Sn 12 8 10 3 0 0
```

```
[[2]]
# A tibble: 6 x 7
  <chr> JulJu AugAu SepSe OctOc NovNo DecDe
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 Average high in °F Av. high Hi 83 80 72 58 41 27
2 Average low in °F Av. low Lo 64 62 52 40 26 12
3 Days with precipitation Days precip.~ 10 10 9 8 8 8
4 Hours of sunshine Hours sun. Sun 343 296 237 193 115 112
5 Av. precipitation in inch Av. precip~ 4.04 4.3 3.08 2.43 1.77 1.16
6 Av. snowfall in inch Snowfall Sn 0 0 0 1 9 12
```

```
[[3]]
# A tibble: 31 x 7
  Day High°F Low°F `Prec/moinch` `Prec/yrinch` `Snow/moinch` `Snow/yrinch`
  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 1 Jan 23.8 8.3 0.04 0.04 0.39 1
2 2 Jan 23.7 8.2 0.08 0.08 0.71 1.8
3 3 Jan 23.6 8.1 0.12 0.12 1.1 2.8
```

```

4 4 Jan    23.5  7.9      0.12      0.12      1.5      3.8
5 5 Jan    23.5  7.8      0.16      0.16      1.81     4.6
6 6 Jan    23.4  7.7      0.2       0.2       2.2      5.6
7 7 Jan    23.4  7.6      0.24     0.24      2.6      6.6
8 8 Jan    23.3  7.5      0.28     0.28     3.11     7.9
9 9 Jan    23.3  7.4      0.28     0.28     3.5      8.9
10 10 Jan   23.3  7.3      0.31     0.31     3.9      9.9
# i 21 more rows

```

```
[[4]]
```

```

# A tibble: 26 x 6
  Day      High°F Low°F Precip.inch Snowinch `Snow d.inch`
  <chr>    <dbl> <dbl> <chr>      <chr>      <dbl>
1 01 Dec    32    19  0.07      1.61        7
2 02 Dec    27    12  0.00      0.00        6
3 03 Dec   37.9  19.9  0.00      0.00        6
4 04 Dec    39   24.1  0.00      0.00        6
5 05 Dec    37   21.9  0.00      0.00        5
6 06 Dec    32   17.1  0.00      0.00        5
7 07 Dec   42.1  21.9  0.00      0.00        5
8 08 Dec    41   30.9  0.00      0.00        5
9 09 Dec    34   -0.9  0.16      2.52        5
10 10 Dec    8.1   -4    T        T        7
# i 16 more rows

```

```
[[5]]
```

```

# A tibble: 9 x 4
  `Dec 19`    `Normal`
  <chr>      <lgl> <chr>
1 "Average high temperature Av. high temp." "29.9 °F" NA "27 °F"
2 "Average low temperature Av. low temp." "14.6 °F" NA "12 °F"
3 "Total precipitation Total precip." "0.39 inch" NA "1.16 inch"
4 "Total snowfall Total snowfall" "6.33 inch" NA "12 inch"
5 "" "" NA ""
6 "Highest max temperature Highest max temp." "44.1 °F" NA "-"
7 "Lowest max temperature Lowest max temp." "8.1 °F" NA "-"
8 "Highest min temperature Highest min temp." "32.0 °F" NA "-"
9 "Lowest min temperature Lowest min temp." "-5.1 °F" NA "-"

```

```
[[6]]
```

```

# A tibble: 10 x 3
  <chr>      <chr>      <lgl>

```

1 Country	United States	NA
2 State	Minnesota	NA
3 County	Hennepin	NA
4 City	Minneapolis	NA
5 Zip code	55401	NA
6 Longitude	-93.27 dec. degr.	NA
7 Latitude	44.98 dec. degr.	NA
8 Altitude - Elevation	840ft	NA
9 ICAO	-	NA
10 IATA	-	NA

[[7]]

```
# A tibble: 6 x 3
  <chr> <chr> <lgl>
1 Local Time 07:23 PM NA
2 Sunrise 07:05 AM NA
3 Sunset 07:32 PM NA
4 Day / Night Day NA
5 Timezone Chicago -6:00 NA
6 Timezone DB America/Chicago NA
```

[[8]]

```
# A tibble: 6 x 2
  <chr> <chr>
1 Annual high temperature 55°F
2 Annual low temperature 37°F
3 Days per year with precip. 112 days
4 Annual hours of sunshine 2607 hours
5 Average annual precip. 30.61 inch
6 Av. annual snowfall 55 inch
```

```
mpls_data1 <- html_table(tables, header = TRUE, fill = TRUE)[[1]]
mpls_data1
```

```
# A tibble: 6 x 7
  <chr> JanJa FebFe MarMa AprAp MayMa JunJu
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 Average high in °F Av. high Hi 24 29 41 58 69 79
2 Average low in °F Av. low Lo 8 13 24 37 49 59
3 Days with precipitation Days precip.~ 8 7 11 9 11 13
```

4	Hours of sunshine	Hours sun.	Sun	140	166	200	231	272	302
5	Av. precipitation in inch	Av. precip~		0.9	0.77	1.89	2.66	3.36	4.25
6	Av. snowfall in inch	Snowfall Sn		12	8	10	3	0	0

```
mpls_data2 <- html_table(tables, header = TRUE, fill = TRUE)[[2]]
mpls_data2
```

```
# A tibble: 6 x 7
```

		JulJu	AugAu	SepSe	OctOc	NovNo	DecDe
<chr>		<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	Average high in °F Av. high Hi	83	80	72	58	41	27
2	Average low in °F Av. low Lo	64	62	52	40	26	12
3	Days with precipitation Days precip.~	10	10	9	8	8	8
4	Hours of sunshine Hours sun. Sun	343	296	237	193	115	112
5	Av. precipitation in inch Av. precip~	4.04	4.3	3.08	2.43	1.77	1.16
6	Av. snowfall in inch Snowfall Sn	0	0	0	1	9	12

Now we wrap the 4 steps above into the `bow` and `scrape` functions from the `polite` package:

```
session <- bow("https://www.usclimatedata.com/climate/minneapolis/minnesota/united-states/us")

result <- scrape(session) |>
  html_nodes(css = "table") |>
  html_table(header = TRUE, fill = TRUE)
mpls_data1 <- result[[1]]
mpls_data2 <- result[[2]]
```

Even after finding the correct tables, there may still be a lot of work to make it tidy!!!

[Pause to Ponder:] What is each line of code doing below?

```
bind_cols(mpls_data1, mpls_data2) |>
  as_tibble() |>
  select(-`...8`) |>
  mutate(`...1` = str_extract(`...1`, "[^ ]+ [^ ]+ [^ ]+")) |>
  pivot_longer(cols = c(`JanJa`: `DecDe`),
               names_to = "month", values_to = "weather") |>
  pivot_wider(names_from = `...1`, values_from = weather) |>
  mutate(month = str_sub(month, 1, 3)) |>
  rename(avg_high = "Average high in",
         avg_low = "Average low in")
```

```

New names:
* `` -> `...1`
* `` -> `...8`

# A tibble: 12 x 7
  month avg_high avg_low `Days with precipitation` `Hours of sunshine`
  <chr>   <dbl>   <dbl>                <dbl>                <dbl>
1 Jan      24      8                      8                    140
2 Feb      29     13                      7                    166
3 Mar      41     24                     11                   200
4 Apr      58     37                      9                   231
5 May      69     49                     11                   272
6 Jun      79     59                     13                   302
7 Jul      83     64                     10                   343
8 Aug      80     62                     10                   296
9 Sep      72     52                      9                   237
10 Oct      58     40                      8                   193
11 Nov      41     26                      8                   115
12 Dec      27     12                      8                   112
# i 2 more variables: `Av. precipitation in` <dbl>, `Av. snowfall in` <dbl>

# Probably want to rename the rest of the variables too!

```

Leaflet mapping example with data in table form

Let's return to our example from 02_maps.qmd where we recreated an [interactive choropleth map](#) of population densities by US state. Recall how that plot was very suspicious? The population density data that came with the state geometries from [our source](#) seemed incorrect.

Let's see if we can use our new web scraping skills to scrape the correct population density data and repeat that plot! Can we go out and find the real statewide population densities, create a tidy data frame, merge that with our state geometry shapefiles, and then regenerate our plot?

A quick wikipedia search yields [this webpage](#) with more reasonable population densities in a nice table format. Let's see if we can grab this data using our 4 steps to **rvesting** data!

```

# check that scraping is allowed (Step 0)
robotstxt::paths_allowed("https://en.wikipedia.org/wiki/List_of_states_and_territories_of_the_United_States")

```

en.wikipedia.org

```
[1] TRUE
```

```
# Step 1: read_html()
pop_dens <- read_html("https://en.wikipedia.org/wiki/List_of_states_and_territories_of_the_U
# 2: html_nodes()
tables <- html_nodes(pop_dens, css = "table")
tables # have to guesstimate which table contains our desired info
```

```
{xml_nodeset (2)}
[1] <table class="wikitable sortable plainrowheaders sticky-header-multi stat ...
[2] <table class="nowraplinks hlist mw-collapsible mw-collapsed navbox-inner" ...
```

```
# 3: html_table()
html_table(tables, header = TRUE, fill = TRUE) # find the right table
```

```
[[1]]
```

```
# A tibble: 61 x 6
```

	Location	Density	Density	Population	`Land area`	`Land area`
	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
1	Location	/mi2	/km2	Population	mi2	km2
2	District of Columbia	11,131	4,297	678,972	61	158
3	New Jersey	1,263	488	9,290,841	7,354	19,047
4	Rhode Island	1,060	409	1,095,962	1,034	2,678
5	Puerto Rico	936	361	3,205,691	3,424	8,868
6	Massachusetts	898	347	7,001,399	7,800	20,202
7	Guam[4]	824	319	172,952	210	543
8	Connecticut	747	288	3,617,176	4,842	12,542
9	U.S. Virgin Islands[4]	737	284	98,750	134	348
10	Maryland	637	246	6,180,253	9,707	25,142

```
# i 51 more rows
```

```
[[2]]
```

```
# A tibble: 11 x 2
```

	.mw-parser-output .navbar{display:inline;font-size:8~1 .mw-parser-output .n~2	
<chr>	<chr>	<chr>
1	"List of states and territories of the United States"	"List of states and t~
2	"Demographics"	"Population\nAfrican ~
3	"Economy"	"Billionaires\nBudget~
4	"Environment"	"Botanical gardens\nC~
5	"Geography"	"Area\nBays\nBeaches~

```

6 "Government" "Agriculture commissi~
7 "Health" "Changes in life expe~
8 "History" "Date of statehood\nN~
9 "Law" "Abortion\nAge of con~
10 "Miscellaneous" "Abbreviations\nAirpo~
11 "Category\n Commons\n Portals" "Category\n Commons\n~
# i abbreviated names:
# 1: `mw-parser-output .navbar{display:inline;font-size:88%;font-weight:normal}.mw-parser~
# 2: `mw-parser-output .navbar{display:inline;font-size:88%;font-weight:normal}.mw-parser~

```

```

density_table <- html_table(tables, header = TRUE, fill = TRUE)[[1]]
density_table

```

```

# A tibble: 61 x 6
  Location          Density Density Population `Land area` `Land area`
  <chr>            <chr>   <chr>   <chr>      <chr>      <chr>
1 Location          /mi2    /km2    Population mi2      km2
2 District of Columbia 11,131  4,297    678,972   61      158
3 New Jersey         1,263   488     9,290,841 7,354    19,047
4 Rhode Island        1,060   409     1,095,962 1,034    2,678
5 Puerto Rico         936     361     3,205,691 3,424    8,868
6 Massachusetts       898     347     7,001,399 7,800    20,202
7 Guam[4]             824     319     172,952   210     543
8 Connecticut         747     288     3,617,176 4,842    12,542
9 U.S. Virgin Islands[4] 737     284     98,750    134     348
10 Maryland           637     246     6,180,253 9,707    25,142
# i 51 more rows

```

```

# Perform Steps 0-3 using the polite package
session <- bow("https://en.wikipedia.org/wiki/List_of_states_and_territories_of_the_United_S~

result <- scrape(session) |>
  html_nodes(css = "table") |>
  html_table(header = TRUE, fill = TRUE)
density_table <- result[[1]]
density_table

```

```

# A tibble: 61 x 6
  Location          Density Density Population `Land area` `Land area`
  <chr>            <chr>   <chr>   <chr>      <chr>      <chr>
1 Location          /mi2    /km2    Population mi2      km2

```

2	District of Columbia	11,131	4,297	678,972	61	158
3	New Jersey	1,263	488	9,290,841	7,354	19,047
4	Rhode Island	1,060	409	1,095,962	1,034	2,678
5	Puerto Rico	936	361	3,205,691	3,424	8,868
6	Massachusetts	898	347	7,001,399	7,800	20,202
7	Guam[4]	824	319	172,952	210	543
8	Connecticut	747	288	3,617,176	4,842	12,542
9	U.S. Virgin Islands[4]	737	284	98,750	134	348
10	Maryland	637	246	6,180,253	9,707	25,142

i 51 more rows

Even after grabbing our table from wikipedia and setting it in a nice tibble format, there is still some cleaning to do before we can merge this with our state geometries:

```
density_data <- density_table |>
  select(1, 2, 4, 5) |>
  filter(!row_number() == 1) |>
  rename(Land_area = `Land area`) |>
  mutate(state_name = str_to_lower(as.character(Location)),
         Density = parse_number(Density),
         Population = parse_number(Population),
         Land_area = parse_number(Land_area)) |>
  select(-Location)
density_data
```

```
# A tibble: 60 x 4
  Density Population Land_area state_name
  <dbl>      <dbl>    <dbl> <chr>
1  11131    678972      61 district of columbia
2   1263   9290841    7354 new jersey
3   1060   1095962    1034 rhode island
4    936   3205691    3424 puerto rico
5    898   7001399    7800 massachusetts
6    824   172952     210 guam[4]
7    747   3617176    4842 connecticut
8    737    98750     134 u.s. virgin islands[4]
9    637   6180253    9707 maryland
10   578    43915     76 american samoa[4]
# i 50 more rows
```

As before, we get core geometry data to draw US states and then we'll make sure we can merge our new density data into the core files.


```
# Get info to draw US states for geom_polygon (connect the lat-long points)
states_polygon <- as_tibble(map_data("state")) |>
  select(region, group, order, lat, long)

# See what the state (region) levels look like in states_polygon
unique(states_polygon$region)
```

[1] "alabama"	"arizona"	"arkansas"
[4] "california"	"colorado"	"connecticut"
[7] "delaware"	"district of columbia"	"florida"
[10] "georgia"	"idaho"	"illinois"
[13] "indiana"	"iowa"	"kansas"
[16] "kentucky"	"louisiana"	"maine"
[19] "maryland"	"massachusetts"	"michigan"
[22] "minnesota"	"mississippi"	"missouri"
[25] "montana"	"nebraska"	"nevada"
[28] "new hampshire"	"new jersey"	"new mexico"
[31] "new york"	"north carolina"	"north dakota"
[34] "ohio"	"oklahoma"	"oregon"
[37] "pennsylvania"	"rhode island"	"south carolina"
[40] "south dakota"	"tennessee"	"texas"
[43] "utah"	"vermont"	"virginia"
[46] "washington"	"west virginia"	"wisconsin"
[49] "wyoming"		

```
# Get info to draw US states for geom_sf and leaflet (simple features
#   object with multipolygon geometry column)
states_sf <- read_sf("https://rstudio.github.io/leaflet/json/us-states.geojson") |>
  select(name, geometry)

# See what the state (name) levels look like in states_sf
unique(states_sf$name)
```

[1] "Alabama"	"Alaska"	"Arizona"
[4] "Arkansas"	"California"	"Colorado"
[7] "Connecticut"	"Delaware"	"District of Columbia"
[10] "Florida"	"Georgia"	"Hawaii"
[13] "Idaho"	"Illinois"	"Indiana"
[16] "Iowa"	"Kansas"	"Kentucky"
[19] "Louisiana"	"Maine"	"Maryland"
[22] "Massachusetts"	"Michigan"	"Minnesota"

[25] "Mississippi"	"Missouri"	"Montana"
[28] "Nebraska"	"Nevada"	"New Hampshire"
[31] "New Jersey"	"New Mexico"	"New York"
[34] "North Carolina"	"North Dakota"	"Ohio"
[37] "Oklahoma"	"Oregon"	"Pennsylvania"
[40] "Rhode Island"	"South Carolina"	"South Dakota"
[43] "Tennessee"	"Texas"	"Utah"
[46] "Vermont"	"Virginia"	"Washington"
[49] "West Virginia"	"Wisconsin"	"Wyoming"
[52] "Puerto Rico"		

```
# See what the state (state_name) levels look like in density_data
unique(density_data$state_name)
```

[1] "district of columbia"	"new jersey"
[3] "rhode island"	"puerto rico"
[5] "massachusetts"	"guam[4]"
[7] "connecticut"	"u.s. virgin islands[4]"
[9] "maryland"	"american samoa[4]"
[11] "delaware"	"florida"
[13] "new york"	"pennsylvania"
[15] "ohio"	"northern mariana islands[4]"
[17] "california"	"illinois"
[19] "hawaii"	"north carolina"
[21] "virginia"	"georgia"
[23] "indiana"	"south carolina"
[25] "michigan"	"tennessee"
[27] "new hampshire"	"washington"
[29] "texas"	"kentucky"
[31] "wisconsin"	"louisiana"
[33] "alabama"	"missouri"
[35] "west virginia"	"minnesota"
[37] "vermont"	"arizona"
[39] "mississippi"	"oklahoma"
[41] "arkansas"	"iowa"
[43] "colorado"	"maine"
[45] "oregon"	"utah"
[47] "kansas"	"nevada"
[49] "nebraska"	"idaho"
[51] "new mexico"	"south dakota"
[53] "north dakota"	"montana"
[55] "wyoming"	"alaska"

```
[57] "contiguous us"          "50 states"
[59] "50 states and dc"       "united states"
```

```
# all lower case plus some extraneous rows
```

```
# Make sure all keys have the same format before joining: all lower case
```

```
states_sf <- states_sf |>
  mutate(name = str_to_lower(name))
```

```
# Now we can merge data sets together for the static and the interactive plots
```

```
# Merge with states_polygon (static)
density_polygon <- states_polygon |>
  left_join(density_data, by = c("region" = "state_name"))
density_polygon
```

```
# A tibble: 15,537 x 8
```

	region	group	order	lat	long	Density	Population	Land_area
	<chr>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	alabama	1	1	30.4	-87.5	101	5108468	50645
2	alabama	1	2	30.4	-87.5	101	5108468	50645
3	alabama	1	3	30.4	-87.5	101	5108468	50645
4	alabama	1	4	30.3	-87.5	101	5108468	50645
5	alabama	1	5	30.3	-87.6	101	5108468	50645
6	alabama	1	6	30.3	-87.6	101	5108468	50645
7	alabama	1	7	30.3	-87.6	101	5108468	50645
8	alabama	1	8	30.3	-87.6	101	5108468	50645
9	alabama	1	9	30.3	-87.7	101	5108468	50645
10	alabama	1	10	30.3	-87.8	101	5108468	50645

```
# i 15,527 more rows
```

```
# Looks like merge worked for 48 contiguous states plus DC
```

```
density_polygon |>
  group_by(region) |>
  summarise(mean = mean(Density)) |>
  print(n = Inf)
```

```
# A tibble: 49 x 2
```

region	mean
<chr>	<dbl>

1	alabama	101
2	arizona	65
3	arkansas	59
4	california	250
5	colorado	57
6	connecticut	747
7	delaware	529
8	district of columbia	11131
9	florida	422
10	georgia	192
11	idaho	24
12	illinois	226
13	indiana	192
14	iowa	57
15	kansas	36
16	kentucky	115
17	louisiana	106
18	maine	45
19	maryland	637
20	massachusetts	898
21	michigan	178
22	minnesota	72
23	mississippi	63
24	missouri	90
25	montana	7.8
26	nebraska	26
27	nevada	29
28	new hampshire	157
29	new jersey	1263
30	new mexico	17
31	new york	415
32	north carolina	223
33	north dakota	11
34	ohio	288
35	oklahoma	59
36	oregon	44
37	pennsylvania	290
38	rhode island	1060
39	south carolina	179
40	south dakota	12
41	tennessee	173
42	texas	117
43	utah	42

```

44 vermont          70
45 virginia         221
46 washington       118
47 west virginia     74
48 wisconsin         109
49 wyoming           6

```

```

# Remove DC since such an outlier
density_polygon <- density_polygon |>
  filter(region != "district of columbia")

# Merge with states_sf (static or interactive)
density_sf <- states_sf |>
  left_join(density_data, by = c("name" = "state_name")) |>
  filter(!(name %in% c("alaska", "hawaii")))

# Looks like merge worked for 48 contiguous states plus DC and PR
class(density_sf)

```

```
[1] "sf"          "tbl_df"      "tbl"         "data.frame"
```

```
print(density_sf, n = Inf)
```

Simple feature collection with 50 features and 4 fields

Geometry type: MULTIPOLYGON

Dimension: XY

Bounding box: xmin: -124.7066 ymin: 17.92956 xmax: -65.6268 ymax: 49.38362

Geodetic CRS: WGS 84

A tibble: 50 x 5


	name	geometry	Density	Population	Land_area
* <chr>		<MULTIPOLYGON [°]>	<dbl>	<dbl>	<dbl>
1	alabama	(((−87.3593 35.00118, −85.~	101	5108468	50645
2	arizona	(((−109.0425 37.00026, −10~	65	7431344	113594
3	arkansas	(((−94.47384 36.50186, −90~	59	3067732	52035
4	california	(((−123.2333 42.00619, −12~	250	38965193	155779
5	colorado	(((−107.9197 41.00391, −10~	57	5877610	103642
6	connecticut	(((−73.05353 42.03905, −71~	747	3617176	4842
7	delaware	(((−75.41409 39.80446, −75~	529	1031890	1949
8	district of columbia	(((−77.03526 38.99387, −76~	11131	678972	61
9	florida	(((−85.49714 30.99754, −85~	422	22610726	53625

10	georgia	(((-83.10919 35.00118, -83~	192	11029227	57513
11	idaho	(((-116.0475 49.00024, -11~	24	1964726	82643
12	illinois	(((-90.63998 42.51006, -88~	226	12549689	55519
13	indiana	(((-85.99006 41.75972, -84~	192	6862199	35826
14	iowa	(((-91.36842 43.50139, -91~	57	3207004	55857
15	kansas	(((-101.906 40.00163, -95.~	36	2940546	81759
16	kentucky	(((-83.90335 38.76931, -83~	115	4526154	39486
17	louisiana	(((-93.60849 33.01853, -91~	106	4573749	43204
18	maine	(((-70.70392 43.05776, -70~	45	1395722	30843
19	maryland	(((-75.99465 37.95325, -76~	637	6180253	9707
20	massachusetts	(((-70.91752 42.88797, -70~	898	7001399	7800
21	michigan	(((-83.45424 41.73234, -84~	178	10037261	56539
22	minnesota	(((-92.0147 46.7054, -92.0~	72	5737915	79627
23	mississippi	(((-88.47111 34.9957, -88.~	63	2939690	46923
24	missouri	(((-91.83396 40.60957, -91~	90	6196156	68742
25	montana	(((-104.0475 49.00024, -10~	7.8	1132812	145546
26	nebraska	(((-103.3246 43.00299, -10~	26	1978379	76824
27	nevada	(((-117.0279 42.00071, -11~	29	3194176	109781
28	new hampshire	(((-71.08183 45.3033, -71.~	157	1402054	8953
29	new jersey	(((-74.23655 41.14083, -73~	1263	9290841	7354
30	new mexico	(((-107.4213 37.00026, -10~	17	2114371	121298
31	new york	(((-73.34381 45.01303, -73~	415	19571216	47126
32	north carolina	(((-80.97866 36.56211, -80~	223	10835491	48618
33	north dakota	(((-97.22874 49.00024, -97~	11	783926	69001
34	ohio	(((-80.5186 41.9788, -80.5~	288	11785935	40861
35	oklahoma	(((-100.0877 37.00026, -94~	59	4053824	68595
36	oregon	(((-123.2113 46.17414, -12~	44	4233358	95988
37	pennsylvania	(((-79.76278 42.25265, -79~	290	12961683	44743
38	rhode island	(((-71.19684 41.67757, -71~	1060	1095962	1034
39	south carolina	(((-82.76414 35.0669, -82.~	179	5373555	30061
40	south dakota	(((-104.0475 45.94411, -96~	12	919318	75811
41	tennessee	(((-88.05487 36.49638, -88~	173	7126489	41235
42	texas	(((-101.8129 36.50186, -10~	117	30503301	261232
43	utah	(((-112.1644 41.99523, -11~	42	3417734	82170
44	vermont	(((-71.50355 45.01303, -71~	70	647464	9217
45	virginia	(((-75.39766 38.0135, -75.~	221	8715698	39490
46	washington	(((-117.0334 49.00024, -11~	118	7812880	66456
47	west virginia	(((-80.5186 40.63695, -80.~	74	1770071	24038
48	wisconsin	(((-90.41543 46.56848, -90~	109	5910955	54158
49	wyoming	(((-109.0808 45.00207, -10~	6	584057	97093
50	puerto rico	(((-66.44834 17.98433, -66~	936	3205691	3424

```
# Remove DC and PR
density_sf <- density_sf |>
  filter(name != "district of columbia" & name != "puerto rico")
```

Numeric variable (static plot):

```
density_polygon |>
  ggplot(mapping = aes(x = long, y = lat, group = group)) +
    geom_polygon(aes(fill = Density), color = "black") +
    labs(fill = "Population density in 2023 \n (people per sq mile)") +
    coord_map() +
    theme_void() +
    scale_fill_viridis()
```

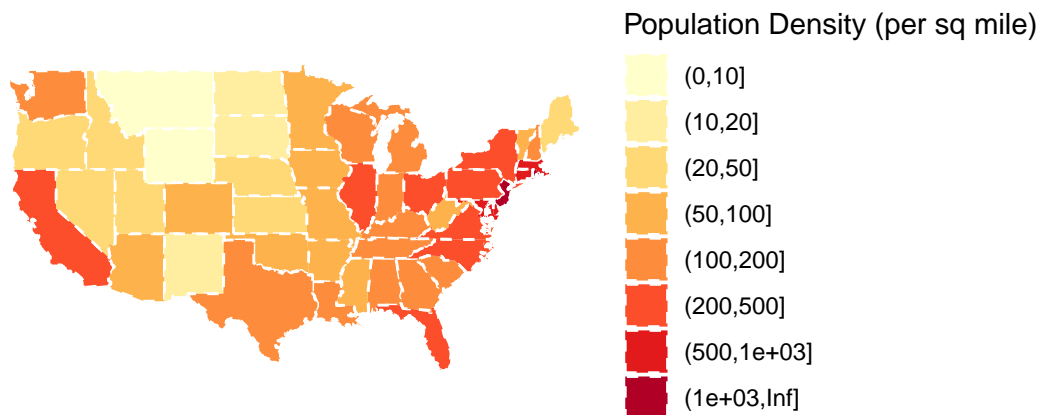


Remember that the original plot classified densities into our own pre-determined bins before plotting - this might look better!

```
density_polygon <- density_polygon |>
  mutate(Density_intervals = cut(Density, n = 8,
    breaks = c(0, 10, 20, 50, 100, 200, 500, 1000, Inf)))

density_polygon |>
```


```
ggplot(mapping = aes(x = long, y = lat, group = group)) +
  geom_polygon(aes(fill = Density_intervals), color = "white",
               linetype = 2) +
  labs(fill = "Population Density (per sq mile)") +
  coord_map() +
  theme_void() +
  scale_fill_brewer(palette = "YlOrRd")
```



We could even create a static plot using `geom_sf()` using `density_sf`:

```
density_sf <- density_sf |>
  mutate(Density_intervals = cut(Density, n = 8,
                                breaks = c(0, 10, 20, 50, 100, 200, 500, 1000, Inf)))

ggplot(data = density_sf) +
  geom_sf(aes(fill = Density_intervals), colour = "white", linetype = 2) +
  theme_void() +
  scale_fill_brewer(palette = "YlOrRd")
```

But... why not make an interactive plot instead?


```
density_sf <- density_sf |>
  mutate(labels = str_c(name, ": ", Density, " people per sq mile in 2023"))

labels <- lapply(density_sf$labels, HTML)
pal <- colorNumeric("YlOrRd", density_sf$Density)

leaflet(density_sf) |>
  setView(-96, 37.8, 4) |>
  addTiles() |>
  addPolygons(
    weight = 2,
    opacity = 1,
    color = ~ pal(density_sf$Density),
    dashArray = "3",
    fillOpacity = 0.7,
    highlightOptions = highlightOptions(
      weight = 5,
      color = "#666",
      dashArray = "",
      fillOpacity = 0.7,
      bringToFront = TRUE),
```

```
label = labels,  
labelOptions = labelOptions(  
    style = list("font-weight" = "normal", padding = "3px 8px"),  
    textsize = "15px",  
    direction = "auto"))
```

file:///C:/Users/charl/AppData/Local/Temp/RtmpWep1Pp/file91bc6389c5/widget91bc5e2842aa.html :



```
# should use addLegend() but not trivial without pre-set bins
```

Here's an interactive plot with our own bins:

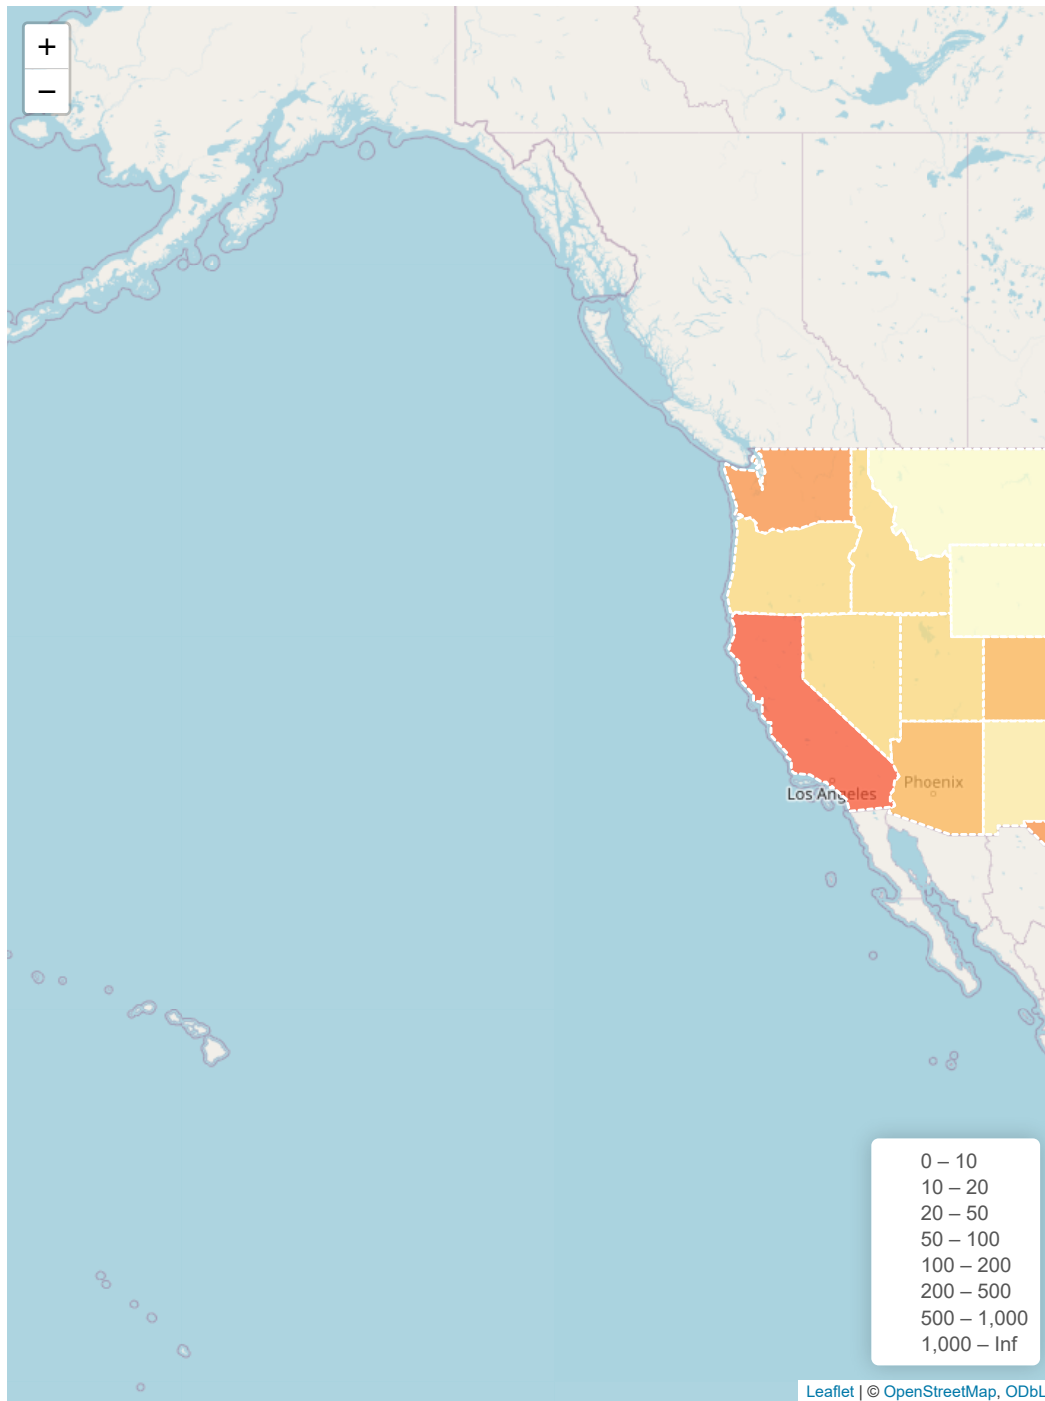
```
# Create our own category bins for population densities
# and assign the yellow-orange-red color palette
bins <- c(0, 10, 20, 50, 100, 200, 500, 1000, Inf)
pal <- colorBin("YlOrRd", domain = density_sf$Density, bins = bins)

# Create labels that pop up when we hover over a state. The labels must
# be part of a list where each entry is tagged as HTML code.
density_sf <- density_sf |>
  mutate(labels = str_c(name, ": ", Density, " people / sq mile"))
labels <- lapply(density_sf$labels, HTML)

# If want more HTML formatting, use these lines instead of those above:
# states <- states |>
#   mutate(labels = glue("<strong>{name}</strong><br/>{density} people /
#   mi<sup>2</sup>"))
# labels <- lapply(states$labels, HTML)

leaflet(density_sf) |>
  setView(-96, 37.8, 4) |>
  addTiles() |>
  addPolygons(
    fillColor = ~pal(Density),
    weight = 2,
    opacity = 1,
    color = "white",
    dashArray = "3",
    fillOpacity = 0.7,
    highlightOptions = highlightOptions(
      weight = 5,
      color = "#666",
      dashArray = "",
      fillOpacity = 0.7,
      bringToFront = TRUE),
    label = labels,
    labelOptions = labelOptions(
      style = list("font-weight" = "normal", padding = "3px 8px"),
      textsize = "15px",
      direction = "auto")) |>
```

```
addLegend(pal = pal, values = ~Density, opacity = 0.7, title = NULL,  
          position = "bottomright")
```



On Your Own

1. Use the `rvest` package and `html_table` to read in the table of data found at the link [here](#) and create a scatterplot of land area versus the 2022 estimated population. I give you some starter code below; fill in the “???” and be sure you can explain what EVERY line of code does and why it’s necessary.

```
#| eval: FALSE
```

```
city_pop <- read_html("https://en.wikipedia.org/wiki/List_of_United_States_cities_by_population")
pop <- html_nodes(???, ???) html_table(pop, header = TRUE, fill = TRUE) # find right
table pop2 <- html_table(pop, header = TRUE, fill = TRUE)[[???]] pop2
```

perform the steps above with the polite package

```
session <- bow("https://en.wikipedia.org/wiki/List_of_United_States_cities_by_population",
force = TRUE)

result <- scrape(session) |> html_nodes(???) |> html_table(header = TRUE, fill = TRUE)
pop2 <- result[[???]] pop2

pop3 <- as_tibble(pop2[,c(1:6,8)]) |> slice(???) |> rename(State = ST, Estimate2023 =
2023estimate, Census = 2020census, Area = 2020 land area, Density = 2020 density)
|> mutate(Estimate2023 = parse_number(Estimate2023), Census = parse_number(Census),
Change = ??? # get rid of % but preserve +/-, Area = parse_number(Area), Density =
parse_number(Density)) |> mutate(City = str_replace(City, "\\.[*$$", "")) pop3
```

pick out unusual points

```
outliers <- pop3 |> filter(Estimate2023 > ??? | Area > ???)
```

This will work if don’t turn variables from chr to dbl, but in that

case notice how axes are just evenly spaced categorical variables

```
ggplot(pop3, aes(x = ???, y = ???)) + geom_point() + geom_smooth() + gg-
pel::geom_label_repel(data = ???, aes(label = ???))
```

2. We would like to create a tibble with 4 years of data (2001-2004) from the Minnesota Wild hockey team. Specifically, we are interested in the “Scoring Regular Season” table from [this webpage](#) and the similar webpages from 2002, 2003, and 2004. Your final tibble should have 6 columns: player, year, age, pos (position), gp (games played), and pts (points).

You should (a) write a function called `hockey_stats` with inputs for team and year to scrape data from the “scoring Regular Season” table, and (b) use iteration techniques to scrape and combine 4 years worth of data. Here are some functions you might consider:

- `row_to_names(row_number = 1)` from the `janitor` package
- `clean_names()` also from the `janitor` package
- `bow()` and `scrape()` from the `polite` package
- `str_c()` from the `stringr` package (for creating urls with user inputs)
- `map2()` and `list_rbind()` for iterating and combining years

Try following these steps:

- 1) Be sure you can find and clean the correct table from the 2021 season.

```
# Step 0: Check that scraping is allowed
robotstxt::paths_allowed("https://www.hockey-reference.com/teams/MIN/2001.html")
```

`www.hockey-reference.com`

```
[1] TRUE
```

```
# Step 1: read_html()
hockey_page <- read_html("https://www.hockey-reference.com/teams/MIN/2001.html")

# Step 2: html_nodes()
tables <- html_nodes(hockey_page, css = "table")
tables # have to guesstimate which table contains our desired info
```

```
{xml_nodeset (6)}
[1] <table class="sortable stats_table" id="team_stats" data-cols-to-freeze=" ...
[2] <table class="sortable stats_table" id="team_stats_adv" data-cols-to-free ...
[3] <table class="sortable stats_table" id="roster" data-cols-to-freeze=",2"> ...
[4] <table class="stats_table sortable per_toggler soc" id="player_stats" dat ...
[5] <table class="stats_table sortable per_toggler soc" id="goalie_stats" dat ...
[6] <table class="stats_table sortable per_toggler soc" id="stats_misc_plus" ...
```



```
# Step 3: html_table()
html_table(tables, header = TRUE, fill = TRUE) # find the right table
```

```
[[1]]
```

```
# A tibble: 2 x 29
```

	Team	AvAge	GP	W	L	T	OL	PTS	`PTS%`	GF	GA	SRS	SOS
	<chr>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<dbl>
1	Minn~	27.4	82	25	39	13	5	68	0.415	168	210	-0.42	0.09
2	Leag~	27.8	82	36	32	10	4	86	0.525	226	226	NA	NA

```
# i 16 more variables: `GF/G` <dbl>, `GA/G` <dbl>, PP <int>, PPO <int>,
# `PP%` <dbl>, PPA <int>, PPOA <int>, `PK%` <dbl>, SH <int>, SHA <int>,
# S <int>, `S%` <dbl>, SA <int>, `SV%` <dbl>, PDO <lgl>, SO <int>
```

```
[[2]]
```

```
# A tibble: 2 x 22
```

	Team	`S%`	`SV%`	PDO	CF	CA	`CF%`	xGF	xGA	aGF	aGA	axDiff	SCF
	<chr>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>
1	Minn~	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
2	Leag~	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA

```
# i 9 more variables: SCA <lgl>, `SCF%` <lgl>, HDF <lgl>, HDA <lgl>,
# `HDF%` <lgl>, HDGF <lgl>, `HDC%` <lgl>, HDGA <lgl>, `HDCO%` <lgl>
```

```
[[3]]
```

```
# A tibble: 38 x 11
```

	No.	Player	Birth	Pos	Age	Ht	Wt	`S/C`	Exp	`Birth Date`	Summary
	<chr>	<chr>	<chr>	<chr>	<int>	<chr>	<int>	<chr>	<chr>	<chr>	<chr>
1	40	Chris A~	ca CA	D	25	6-0	205	L/-	R	June 26, 19~	0 G, 0~
2	45	Peter B~	cs CS	RW	27	6-0	185	R/-	R	September 5~	4 G, 2~
3	3	Ladislav~	cs CS	D	25	6-2	190	L/-	1	March 24, 1~	2 G, 5~
4	31	Zac Bie~	ca CA	G	24	6-5	205	-/L	3	September 1~	0-1-0,~
5	36	Sylvain~	ca CA	LW	26	6-2	215	L/-	3	May 21, 1974	3 G, 2~
6	5	Brad Bo~	ca CA	D	28	6-1	205	L/-	3	May 5, 1972	0 G, 1~
7	32	Brian B~	us US	LW	27	5-10	186	L/-	1	November 28~	0 G, 0~
8	15	J.J. Da~	ca CA	D	35	5-10	192	L/-	15	October 12,~	0 G, 0~
9	34	Jim Dowd	us US	C	32	6-0	180	R/-	9	December 25~	7 G, 2~
10	11	Pascal ~	ca CA	LW	21	6-1	205	L/-	R	April 7, 19~	1 G, 0~

```
# i 28 more rows
```

```
[[4]]
```

```
# A tibble: 40 x 22
```

						Scoring	Scoring	Scoring			Goals	Goals
	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>

1 Rk	Play~	Age	Pos	GP	G	A	PTS	+/-	PIM	EVG	PPG
2 1	Scot~	31	RW	58	11	28	39	6	45	7	2
3 2	Mari~	18	LW	71	18	18	36	-6	32	12	6
4 3	Lubo~	32	D	80	11	23	34	-8	52	7	4
5 4	Wes ~	30	C	82	18	12	30	-8	37	11	0
6 5	Fili~	24	D	75	9	21	30	-6	28	5	4
7 6	Darb~	28	LW	72	18	11	29	1	36	14	3
8 7	Jim ~	32	C	68	7	22	29	-6	80	7	0
9 8	Antt~	27	LW	82	12	16	28	-7	24	10	0
10 9	Stac~	26	C	76	7	20	27	3	20	6	1

i 30 more rows

i 10 more variables: Goals <chr>, Goals <chr>, Assists <chr>, Assists <chr>,
 # Assists <chr>, Shots <chr>, Shots <chr>, `Ice Time` <chr>,
 # `Ice Time` <chr>, `` <chr>

[[5]]

A tibble: 6 x 23

				`Goalie Stats`	`Goalie Stats`	`Goalie Stats`	`Goalie Stats`
	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
1 "Rk"	Player	"Age"	GP	W	L	T/O	
2 "1"	Jamie~	"29"	38	5	23	9	
3 "2"	Manny~	"26"	42	19	17	4	
4 "3"	Derek~	"21"	4	1	3	0	
5 "4"	Zac B~	"24"	1	0	1	0	
6 ""	Team ~	"	82	25	44	13	

i 16 more variables: `Goalie Stats` <chr>, `Goalie Stats` <chr>,
 # `Goalie Stats` <chr>, `Goalie Stats` <chr>, `Goalie Stats` <chr>,
 # `Goalie Stats` <chr>, `Goalie Stats` <chr>, `Goalie Stats` <chr>,
 # `Goalie Stats` <chr>, `Goalie Stats` <chr>, `Goalie Stats` <chr>,
 # Scoring <chr>, Scoring <chr>, Scoring <chr>, `` <chr>, `` <chr>

[[6]]

A tibble: 35 x 18

						Adjusted	Adjusted	Adjusted	Adjusted
	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
1 Rk	Player	Age	Pos	GP	G	A	PTS	GC	
2 1	Scott Pellerin	31	RW	58	12	30	42	14.7	
3 2	Marián Gáborík	18	LW	71	20	19	39	16.1	
4 3	Lubomír Sekeráš	32	D	80	12	25	37	13.4	
5 4	Wes Walz	30	C	82	20	13	33	14.5	
6 5	Filip Kuba	24	D	75	10	22	32	11.5	
7 6	Jim Dowd	32	C	68	8	23	31	10.6	
8 7	Darby Hendrickson	28	LW	72	20	12	32	14.2	

```

 9 8      Antti Laaksonen    27    LW    82    13        17        30        11.7
10 9      Stacy Roest       26     C     76     8        21        29        10.1
# i 25 more rows
# i 9 more variables: `Plus/Minus` <chr>, `Plus/Minus` <chr>,
#   `Plus/Minus` <chr>, `Plus/Minus` <chr>, `Plus/Minus` <chr>,
#   `Point Shares` <chr>, `Point Shares` <chr>, `Point Shares` <chr>, `` <chr>

```

```

hockey_table <- html_table(tables, header = TRUE, fill = TRUE)[[1]]
hockey_table

```

```

# A tibble: 2 x 29
  Team AvAge GP W L T OL PTS `PTS%` GF GA SRS SOS
  <chr> <dbl> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <dbl>
1 Minn~ 27.4 82 25 39 13 5 68 0.415 168 210 -0.42 0.09
2 Leag~ 27.8 82 36 32 10 4 86 0.525 226 226 NA NA
# i 16 more variables: `GF/G` <dbl>, `GA/G` <dbl>, PP <int>, PPO <int>,
#   `PP%` <dbl>, PPA <int>, PPOA <int>, `PK%` <dbl>, SH <int>, SHA <int>,
#   S <int>, `S%` <dbl>, SA <int>, `SV%` <dbl>, PDO <lgl>, SO <int>

```

2) Organize your `rvest` code from (1) into functions from the `polite` package.

```

session <- bow("https://www.hockey-reference.com/teams/MIN/2001.html", force = TRUE)

result <- scrape(session) |>
  html_nodes(css = "table") |>
  html_table(header = TRUE, fill = TRUE)

```

No encoding supplied: defaulting to UTF-8.

```

hockey_table <- result[[1]]
hockey_table

```

```

# A tibble: 2 x 29
  Team AvAge GP W L T OL PTS `PTS%` GF GA SRS SOS
  <chr> <dbl> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <dbl>
1 Minn~ 27.4 82 25 39 13 5 68 0.415 168 210 -0.42 0.09
2 Leag~ 27.8 82 36 32 10 4 86 0.525 226 226 NA NA
# i 16 more variables: `GF/G` <dbl>, `GA/G` <dbl>, PP <int>, PPO <int>,
#   `PP%` <dbl>, PPA <int>, PPOA <int>, `PK%` <dbl>, SH <int>, SHA <int>,
#   S <int>, `S%` <dbl>, SA <int>, `SV%` <dbl>, PDO <lgl>, SO <int>

```

- 3) Place the code from (2) into a function where the user can input a team and year. You would then adjust the url accordingly and produce a clean table for the user.

```
hockey_stats <- function(team, year){
  base_front_url <- "https://www.hockey-reference.com/teams/"
  url <- str_c(base_front_url, team, "/", year, ".html")
  session <- bow(url, force = TRUE)

  result <- scrape(session) |>
    html_nodes(css = "table") |>
    html_table(header = TRUE, fill = TRUE)
  hockey_table <- result[[1]]
  hockey_table
}
```

```
hockey_stats("MIN", "2001")
```

No encoding supplied: defaulting to UTF-8.

A tibble: 2 x 29

	Team	AvAge	GP	W	L	T	OL	PTS	`PTS%`	GF	GA	SRS	SOS
	<chr>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<dbl>
1	Minn~	27.4	82	25	39	13	5	68	0.415	168	210	-0.42	0.09
2	Leag~	27.8	82	36	32	10	4	86	0.525	226	226	NA	NA

i 16 more variables: `GF/G` <dbl>, `GA/G` <dbl>, PP <int>, PPO <int>,
`PP%` <dbl>, PPA <int>, PPOA <int>, `PK%` <dbl>, SH <int>, SHA <int>,
S <int>, `S%` <dbl>, SA <int>, `SV%` <dbl>, PDO <lgl>, SO <int>

- 4) Use map2 and list_rbind to build one data set containing Minnesota Wild data from 2001-2004.

```
specific_years <- c("2001","2002","2003","2004")
mn_hockey_data <- map2("MIN", specific_years, hockey_stats) |>
  list_rbind()
```

No encoding supplied: defaulting to UTF-8.

No encoding supplied: defaulting to UTF-8.

No encoding supplied: defaulting to UTF-8.

No encoding supplied: defaulting to UTF-8.

You can download this .qmd file from [here](#). Just hit the Download Raw File button.

Credit to Brianna Heggseth and Leslie Myint from Macalester College for a few of these descriptions and examples.

Using rvest for web scraping

Please see `08_table_scraping.qmd` for a preview of web scraping techniques when no API exists, along with ethical considerations when scraping data. In this file, we will turn to scenarios when the webpage contains data of interest, but it is not already in table form.

Recall the four steps to scraping data with functions in the `rvest` library:

0. `robotstxt::paths_allowed()` Check if the website allows scraping, and then make sure we scrape “politely”
1. `read_html()`. Input the URL containing the data and turn the html code into an XML file (another markup format that’s easier to work with).
2. `html_nodes()`. Extract specific nodes from the XML file by using the CSS path that leads to the content of interest. (use `css=“table”` for tables.)
3. `html_text()`. Extract content of interest from nodes. Might also use `html_table()` etc.

More scraping ethics

`robots.txt`

`robots.txt` is a file that some websites will publish to clarify what can and cannot be scraped and other constraints about scraping. When a website publishes this file, this we need to comply with the information in it for moral and legal reasons.

We will look through the information in [this tutorial](#) and apply this to the [NIH robots.txt file](#).

From our investigation of the NIH `robots.txt`, we learn:

- `User-agent: *`: Anyone is allowed to scrape
- `Crawl-delay: 2`: Need to wait 2 seconds between each page scraped
- `No Visit-time` entry: no restrictions on time of day that scraping is allowed
- `No Request-rate` entry: no restrictions on simultaneous requests
- No mention of `?page=`, `news-events`, `news-releases`, or `https://science.education.nih.gov/` in the `Disallow` sections. (This is what we want to scrape today.)

`robotstxt` package

We can also use functions from the [robotstxt package](#), which was built to download and parse `robots.txt` files ([more info](#)). Specifically, the `paths_allowed()` function can check if a bot has permission to access certain pages.

A timeout to preview some technical ideas

HTML structure

HTML (hypertext markup language) is the formatting language used to create webpages. We can see the core parts of HTML from the [rvest vignette](#).

Finding CSS Selectors

In order to gather information from a webpage, we must learn the language used to identify patterns of specific information. For example, on the [NIH News Releases page](#), we can see that the data is represented in a consistent pattern of image + title + abstract.

We will identify data in a web page using a pattern matching language called [CSS Selectors](#) that can refer to specific patterns in HTML, the language used to write web pages.

For example:

- Selecting by tag:
 - "a" selects all hyperlinks in a webpage ("a" represents "anchor" links in HTML)
 - "p" selects all paragraph elements
- Selecting by ID and class:
 - ".description" selects all elements with **class** equal to "description"
 - * The . at the beginning is what signifies **class** selection.
 - * This is one of the most common CSS selectors for scraping because in HTML, the **class** attribute is extremely commonly used to format webpage elements. (Any number of HTML elements can have the same **class**, which is not true for the **id** attribute.)
 - "#mainTitle" selects the **SINGLE** element with **id** equal to "mainTitle"
 - * The # at the beginning is what signifies **id** selection.

```
<p class="title">Title of resource 1</p>
<p class="description">Description of resource 1</p>

<p class="title">Title of resource 2</p>
<p class="description">Description of resource 2</p>
```

Warning: Websites change often! So if you are going to scrape a lot of data, it is probably worthwhile to save and date a copy of the website. Otherwise, you may return after some time and your scraping code will include all of the wrong CSS selectors.

SelectorGadget

Although you can [learn how to use CSS Selectors by hand](#), we will use a shortcut by installing the [Selector Gadget](#) tool.

- There is a version available for Chrome—add it to Chrome via the [Chome Web Store](#).
 - Make sure to pin the extension to the menu bar. (Click the 3 dots > Extensions > Manage extensions. Click the “Details” button under SelectorGadget and toggle the “Pin to toolbar” option.)
- There is also a version that can be saved as a bookmark in the browser—see [here](#).

You might watch the Selector Gadget [tutorial video](#).

Case Study: NIH News Releases

Our goal is to build a data frame with the article title, publication date, and abstract text for the 50 most recent NIH news releases.

Head over to the [NIH News Releases page](#). Click the Selector Gadget extension icon or bookmark button. As you mouse over the webpage, different parts will be highlighted in orange. Click on the title (but not the live link portion!) of the first news release. You’ll notice that the Selector Gadget information in the lower right describes what you clicked on. (If SelectorGadget ever highlights too much in green, you can click on portions that you do not want to turn them red.)

Scroll through the page to verify that only the information you intend (the description paragraph) is selected. The selector panel shows the CSS selector (`.teaser-title`) and the number of matches for that CSS selector (10). (You may have to be careful with your clicking—there are two overlapping boxes, and clicking on the link of the title can lead to the CSS selector of “a”.)

[Pause to Ponder:] Repeat the process above to find the correct selectors for the following fields. Make sure that each matches 10 results:

- The publication date
`.date-display-single`
- The article abstract paragraph (which will also include the publication date)
`.teaser-description`

Retrieving Data Using rvest and CSS Selectors

Now that we have identified CSS selectors for the information we need, let's fetch the data using the `rvest` package similarly to our approach in `08_table_scraping.qmd`.

```
# check that scraping is allowed (Step 0)
robotstxt::paths_allowed("https://www.nih.gov/news-events/news-releases")
```

```
www.nih.gov
```

```
[1] TRUE
```

```
# Step 1: Download the HTML and turn it into an XML file with read_html()
nih <- read_html("https://www.nih.gov/news-events/news-releases")
```

Finding the exact node (e.g. “teaser-title”) is the tricky part. Among all the html code used to produce a webpage, where do you go to grab the content of interest? This is where SelectorGadget comes to the rescue!

```
# Step 2: Extract specific nodes with html_nodes()
title_temp <- html_nodes(nih, ".teaser-title")
title_temp
```

```
{xml_nodeset (10)}
[1] <h4 class="teaser-title"><a href="/news-events/news-releases/nih-researc ...
[2] <h4 class="teaser-title"><a href="/news-events/news-releases/study-illum ...
[3] <h4 class="teaser-title"><a href="/news-events/news-releases/nih-funded- ...
[4] <h4 class="teaser-title"><a href="/news-events/news-releases/surgery-kid ...
[5] <h4 class="teaser-title"><a href="/news-events/news-releases/nih-sponsor ...
[6] <h4 class="teaser-title"><a href="/news-events/news-releases/topical-ste ...
[7] <h4 class="teaser-title"><a href="/news-events/news-releases/tecovirimat ...
[8] <h4 class="teaser-title"><a href="/news-events/news-releases/nih-central ...
[9] <h4 class="teaser-title"><a href="/news-events/news-releases/nih-funded- ...
[10] <h4 class="teaser-title"><a href="/news-events/news-releases/longer-brea ...
```

```
# Step 3: Extract content from nodes with html_text(), html_name(),
#       html_attrs(), html_children(), html_table(), etc.
# Usually will still need to do some stringr adjustments
title_vec <- html_text(title_temp)
title_vec
```



```
[1] "NIH researchers develop eye drops that slow vision loss in animals"
[2] "Study illuminates the structural features of memory formation at cellular and subcellular levels"
[3] "NIH-funded study identifies potential new stroke treatment"
[4] "Surgery in kids with mild sleep-disordered breathing tied to fewer doctor visits, meds"
[5] "NIH-sponsored trial of Lassa vaccine opens"
[6] "Topical steroid withdrawal diagnostic criteria defined by NIH researchers"
[7] "Tecovirimat is safe but ineffective as treatment for clade II mpox"
[8] "NIH centralizes peer review to improve efficiency and strengthen integrity "
[9] "NIH-funded research team engineers new drug targeting pain sensation pathway"
[10] "Longer breastfeeding linked to blood-pressure lowering effects of certain infant gut bacteria"
```

You can also write this altogether with a pipe:

```
robotstxt::paths_allowed("https://www.nih.gov/news-events/news-releases")
```

```
www.nih.gov
```

```
[1] TRUE
```

```
read_html("https://www.nih.gov/news-events/news-releases") |>
  html_nodes(".teaser-title") |>
  html_text()
```

```
[1] "NIH researchers develop eye drops that slow vision loss in animals"
[2] "Study illuminates the structural features of memory formation at cellular and subcellular levels"
[3] "NIH-funded study identifies potential new stroke treatment"
[4] "Surgery in kids with mild sleep-disordered breathing tied to fewer doctor visits, meds"
[5] "NIH-sponsored trial of Lassa vaccine opens"
[6] "Topical steroid withdrawal diagnostic criteria defined by NIH researchers"
[7] "Tecovirimat is safe but ineffective as treatment for clade II mpox"
[8] "NIH centralizes peer review to improve efficiency and strengthen integrity "
[9] "NIH-funded research team engineers new drug targeting pain sensation pathway"
[10] "Longer breastfeeding linked to blood-pressure lowering effects of certain infant gut bacteria"
```

And finally we wrap the 4 steps above into the `bow` and `scrape` functions from the `polite` package:

```
session <- bow("https://www.nih.gov/news-events/news-releases", force = TRUE)

nih_title <- scrape(session) |>
  html_nodes(".teaser-title") |>
  html_text()
nih_title
```

```
[1] "NIH researchers develop eye drops that slow vision loss in animals"
[2] "Study illuminates the structural features of memory formation at cellular and subcellular levels"
[3] "NIH-funded study identifies potential new stroke treatment"
[4] "Surgery in kids with mild sleep-disordered breathing tied to fewer doctor visits, medications"
[5] "NIH-sponsored trial of Lassa vaccine opens"
[6] "Topical steroid withdrawal diagnostic criteria defined by NIH researchers"
[7] "Tecovirimat is safe but ineffective as treatment for clade II mpox"
[8] "NIH centralizes peer review to improve efficiency and strengthen integrity"
[9] "NIH-funded research team engineers new drug targeting pain sensation pathway"
[10] "Longer breastfeeding linked to blood-pressure lowering effects of certain infant gut bacteria"
```

Putting multiple columns of data together.

Now repeat the process above to extract the publication date and the abstract.

```
nih_pubdate <- scrape(session) |>
  html_nodes(".date-display-single") |>
  html_text()
nih_pubdate
```

```
[1] "March 21, 2025" "March 20, 2025" "March 17, 2025" "March 17, 2025"
[5] "March 17, 2025" "March 14, 2025" "March 12, 2025" "March 6, 2025"
[9] "March 5, 2025"  "March 4, 2025"
```

```
nih_description <- scrape(session) |>
  html_nodes(".teaser-description") |>
  html_text()
nih_description
```

```
[1] "March 21, 2025 - \n          Treatment shows potential to slow the progression of L"
[2] "March 20, 2025 - \n          NIH-funded study uses cutting-edge imaging techniques"
[3] "March 17, 2025 - \n          Preclinical study in rodents suggests that uric acid "
```

```

[4] "March 17, 2025 - \n      NIH-funded study supports use of adenotonsillectomy in
[5] "March 17, 2025 - \n      Lassa fever is a viral hemorrhagic disease that can be
[6] "March 14, 2025 - \n      Criteria may help guide treatment of dermatitis. "
[7] "March 12, 2025 - \n      NIH-sponsored trial data offer further evidence to he
[8] "March 6, 2025 - \n      The proposed approach is expected to save more than $6
[9] "March 5, 2025 - \n      Study of CB1 receptor has implications for chronic pain
[10] "March 4, 2025 - \n      Nursing for at least six months may spur beneficial gut

```

Combine these extracted variables into a single tibble. Make sure the variables are formatted correctly - e.g. `pubdate` has date type, `description` does not contain the `pubdate`, etc.

```

# use tibble() to put multiple columns together into a tibble
nih_top10 <- tibble(title = nih_title,
                    pubdate = nih_pubdate,
                    description = nih_description)
nih_top10

```

A tibble: 10 x 3

	title	pubdate	description
	<chr>	<chr>	<chr>
1	"NIH researchers develop eye drops that slow vision loss~	March ~	"March 21,~
2	"Study illuminates the structural features of memory for~	March ~	"March 20,~
3	"NIH-funded study identifies potential new stroke treatm~	March ~	"March 17,~
4	"Surgery in kids with mild sleep-disordered breathing ti~	March ~	"March 17,~
5	"NIH-sponsored trial of Lassa vaccine opens"	March ~	"March 17,~
6	"Topical steroid withdrawal diagnostic criteria defined ~	March ~	"March 14,~
7	"Tecovirimat is safe but ineffective as treatment for cl~	March ~	"March 12,~
8	"NIH centralizes peer review to improve efficiency and s~	March ~	"March 6, ~
9	"NIH-funded research team engineers new drug targeting p~	March ~	"March 5, ~
10	"Longer breastfeeding linked to blood-pressure lowering ~	March ~	"March 4, ~

```

# now clean the data
nih_top10 <- nih_top10 |>
  mutate(pubdate = mdy(pubdate),
         description = str_trim(str_replace(description, ".*\\n", "")))
nih_top10

```

A tibble: 10 x 3

	title	pubdate	description
	<chr>	<date>	<chr>
1	"NIH researchers develop eye drops that slow vision l~	2025-03-21	Treatment ~

```

2 "Study illuminates the structural features of memory ~ 2025-03-20 NIH-funded~
3 "NIH-funded study identifies potential new stroke tre~ 2025-03-17 Preclinica~
4 "Surgery in kids with mild sleep-disordered breathing~ 2025-03-17 NIH-funded~
5 "NIH-sponsored trial of Lassa vaccine opens"          2025-03-17 Lassa feve~
6 "Topical steroid withdrawal diagnostic criteria defin~ 2025-03-14 Criteria m~
7 "Tecovirimat is safe but ineffective as treatment for~ 2025-03-12 NIH-sponso~
8 "NIH centralizes peer review to improve efficiency an~ 2025-03-06 The propos~
9 "NIH-funded research team engineers new drug targetin~ 2025-03-05 Study of C~
10 "Longer breastfeeding linked to blood-pressure loweri~ 2025-03-04 Nursing fo~

```

NOW - continue this process to build a tibble with the most recent 50 NIH news releases, which will require that you iterate over 5 webpages! You should write at least one function, and you will need iteration—use both a `for` loop and appropriate `map_()` functions from `purrr`. Some additional hints:

- Mouse over the page buttons at the very bottom of the news home page to see what the URLs look like.
- Include `Sys.sleep(2)` in your function to respect the `Crawl-delay: 2` in the NIH `robots.txt` file.
- Recall that `bind_rows()` from `dplyr` takes a list of data frames and stacks them on top of each other.

[Pause to Ponder:] Create a function to scrape a single NIH press release page by filling missing pieces labeled ???:

```

# Helper function to reduce html_nodes() |> html_text() code duplication
get_text_from_page <- function(page, css_selector) {
  page |>
    html_nodes(css_selector) |>
    html_text()
}

# Main function to scrape and tidy desired attributes
scrape_page <- function(url) {
  Sys.sleep(2)
  page <- read_html(url)
  article_titles <- get_text_from_page(page, ".teaser-title")
  article_dates <- get_text_from_page(page, ".date-display-single")
  article_dates <- mdy(article_dates)
  article_description <- get_text_from_page(page, ".teaser-description")
  article_description <- str_trim(str_replace(article_description,
                                                ".*\\n",
                                                ""))
}

```

```

    )

    tibble(
      title = article_titles,
      dates = article_dates,
      description = article_description
    )
  }

scrape_page("https://www.nih.gov/news-events/news-releases")

```

```

# A tibble: 10 x 3
  title                                dates      description
  <chr>                                <date>     <chr>
1 "NIH researchers develop eye drops that slow vision l~ 2025-03-21 Treatment ~
2 "Study illuminates the structural features of memory ~ 2025-03-20 NIH-funded~
3 "NIH-funded study identifies potential new stroke tre~ 2025-03-17 Preclinica~
4 "Surgery in kids with mild sleep-disordered breathing~ 2025-03-17 NIH-funded~
5 "NIH-sponsored trial of Lassa vaccine opens"          2025-03-17 Lassa feve~
6 "Topical steroid withdrawal diagnostic criteria defin~ 2025-03-14 Criteria m~
7 "Tecovirimat is safe but ineffective as treatment for~ 2025-03-12 NIH-sponso~
8 "NIH centralizes peer review to improve efficiency an~ 2025-03-06 The propos~
9 "NIH-funded research team engineers new drug targetin~ 2025-03-05 Study of C~
10 "Longer breastfeeding linked to blood-pressure loweri~ 2025-03-04 Nursing fo~

```

[Pause to Ponder:] Use a for loop over the first 5 pages:

```

pages <- vector("list", length = 6)
pos <- 0

for (i in 2025:2024) {
  for (j in 0:2) {
    pos <- pos + 1
    url <- str_c("https://www.nih.gov/news-events/news-releases?", i,
                 "&page=", j, "&1=")
    pages[[pos]] <- scrape_page(url)
  }
}

df_articles <- bind_rows(pages)
head(df_articles)

```

```
# A tibble: 6 x 3
```

	title <chr>	dates <date>	description <chr>
1	NIH researchers develop eye drops that slow vision los~	2025-03-21	Treatment ~
2	Study illuminates the structural features of memory fo~	2025-03-20	NIH-funded~
3	NIH-funded study identifies potential new stroke treat~	2025-03-17	Preclinica~
4	Surgery in kids with mild sleep-disordered breathing t~	2025-03-17	NIH-funded~
5	NIH-sponsored trial of Lassa vaccine opens	2025-03-17	Lassa feve~
6	Topical steroid withdrawal diagnostic criteria defined~	2025-03-14	Criteria m~

[Pause to Ponder:] Use map functions in the purrr package:

```
library(purrr)

base_url <- "https://www.nih.gov/news-events/news-releases?page="
urls_all_pages <- str_c(base_url, seq(0,5))

pages2 <- purrr::map(urls_all_pages, scrape_page)
df_articles2 <- bind_rows(pages2)
head(df_articles2)
```

```
# A tibble: 6 x 3
```

	title <chr>	dates <date>	description <chr>
1	NIH researchers develop eye drops that slow vision los~	2025-03-21	Treatment ~
2	Study illuminates the structural features of memory fo~	2025-03-20	NIH-funded~
3	NIH-funded study identifies potential new stroke treat~	2025-03-17	Preclinica~
4	Surgery in kids with mild sleep-disordered breathing t~	2025-03-17	NIH-funded~
5	NIH-sponsored trial of Lassa vaccine opens	2025-03-17	Lassa feve~
6	Topical steroid withdrawal diagnostic criteria defined~	2025-03-14	Criteria m~